

# Principles of Computer Architecture

---

CSE 240A

Fall 2024

Hadi Esmaeilzadeh

[hadi@ucsd.edu](mailto:hadi@ucsd.edu)

University of California, San Diego



# How to Summarize Performance

---

	Computer A	Computer B	Computer C
Program 1	1	10	20
Program 2	1000	100	20
Total time	1001	110	40

*Which machine is fastest?*

# How to Summarize Performance

---

- Arithmetic Mean  $\frac{1}{n} \sum_{i=1}^n Time_i$
- Weighted Arithmetic Mean  $\sum_{i=1}^n Time_i * Weight_i$   
where the sum of the weights is 1.
- Geometric Mean  $\sqrt[n]{\prod_{i=1}^n ExecutionTimeRatio_i} = \frac{\sqrt[n]{\prod_{i=1}^n ExecutionTime_i}}{ExecutionTime_{base}}$

# Summarizing Performance

Machines:	<u>A</u>	<u>B</u>
Program 1	1	10
Program 2	1000	100

Arith M: Speedup (A/B) =  $(10 / 1 + 100 / 1000) / 2 = 5.05$   
 Arith M: Speedup (A/B) =  $(10 + 100) / (1 + 1000) = 0.10989$

Geo M: Speedup (A/B) =  $\sqrt{\sqrt{10 / 1} * \sqrt{100 / 1000}} = 1$   
 Geo M: Speedup (A/B) =  $\sqrt{10 * 100} / \sqrt{1 * 1000} = 1$

	<u>Set(1)</u>	<u>Set(2)</u>	<u>Set(3)</u>
$W_1$	.5	.909	.999
$W_2$	.5	.091	.001
Arith M/Set(1)		500.5	55
Arith M/Set(2)		91.82	18.18
Arith M/Set(3)		2	10.09
Geo M		31.6	31.6

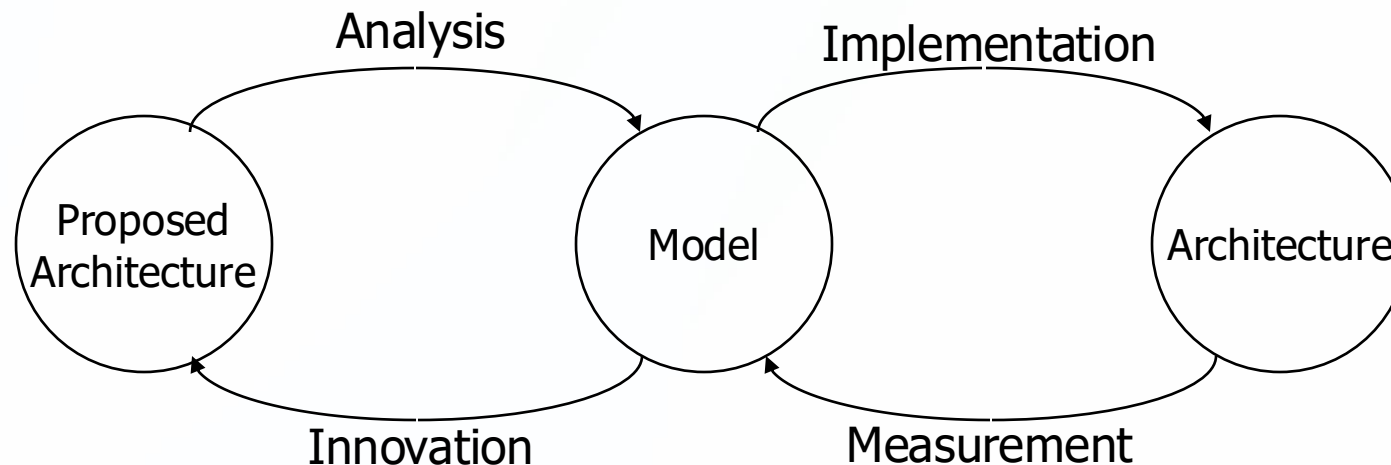
# Summarizing Performance

---

- Even the unweighted arithmetic mean implies a weighting
- Ratios of geometric means never change (regardless of which machine is used as the base), and always give equal weight to all benchmarks
- To give unequal weight requires weighted arithmetic mean

# Analyzing Performance

- That was all about measuring performance. What tools do we use to analyze (predict) performance in the absence of something to measure?
  - models, equations, queueing theory, mean value analysis, instruction-level simulation, gate-level simulation, ...



Because we would like to understand the effects of changes we make to the architecture, we mode/measure performance

# Speedup (due to architectural change)

---

- Speedup is just relative performance on the same machine with something changed.
- From before, then:

$$\text{Speedup} = \text{Relative Performance} = \frac{\text{ET for entire task without change}}{\text{ET for entire task with change}}$$

# Speedup (due to architectural change)

---

- Speedup is just relative performance on the same machine with something changed.
- From before, then:

$$\text{Speedup} = \text{Relative Performance} = \frac{\text{ET for entire task without change}}{\text{ET for entire task with change}}$$

Suppose the change only affects part of execution time...



# Amdahl's Law

The impact of a performance improvement is limited by the percent of execution time affected by the improvement

$$\text{Execution time after improvement} = \frac{\text{Execution Time Affected}}{\text{Amount of Improvement}} + \text{Execution Time Unaffected}$$

# Amdahl's Law

The impact of a performance improvement is limited by the percent of execution time affected by the improvement

$$\text{Execution time after improvement} = \frac{\text{Execution Time Affected}}{\text{Amount of Improvement}} + \text{Execution Time Unaffected}$$

Make the common case fast!!

**TAKE AWAY: Focus your optimizations on where it has the most impact**

# Attendance

---

Selection	Answer
Present	A
Present	B
Present	C
Present	B
Present	D

# Example

- Program A runs for 20 seconds, but 5 seconds of that time is just waiting for memory. If we double the speed of the memory subsystem, what is the speedup?

Selection	Speedup
A	2.00
B	1.50
C	1.28
D	1.14
E	None of the above


# Amdahl's Law in the Era of Multicores



Thread Level Parallelism

# Amdahl's Law in the Era of Multicores

Speedup




1.0

$1/.55 = 1.82$

Thread Level Parallelism

# Amdahl's Law in the Era of Multicores


Speedup



Thread Level Parallelism

# Amdahl's Law in the Era of Multicores

Speedup





# Is there any other way besides TLP?

- Speedup a single application through TLP runs into problems when using multicore because of diminishing returns and dominance of serial execution.


- What if there was a time constraint?

Human perception  
(graphics)

Earthquake prediction

Weather prediction

# Scale Up (Data Level Parallelism instead of TLP)



# Beyond Time...

---

- $\text{Energy} = \text{Power} * \text{Time}$
- $\text{Joules} = \text{Watts} * \text{Seconds}$

# Beyond Time...

---

- Energy = power \* time
- Joules = Watts \* seconds
- If you can halve execution time of a program but to do so, you raise the average power by 50%. Did you save energy?

Selection	Answer
A	Yes
B	No

# Beyond Time...

---

- Energy = power \* time
- Joules = Watts \* seconds
- If you can halve execution time of a program but to do so, you raise the average power by 25%. Did you save energy?
- We'll touch on energy concerns in 240A a bit, but 240C (and 240B, to some extent) will discuss this.

# Back to CPUs, what is Time?

---

# What is Time?

---

$$\begin{aligned}\text{CPU Execution Time} &= \text{CPU clock cycles} * \text{Clock cycle time} \\ &= \text{CPU clock cycles} / \text{Clock rate}\end{aligned}$$

Every conventional processor has a clock with an associated clock cycle time or clock rate.

Every program runs in an integral number of clock cycles.

GHz = billions of cycles/second

X GHz = 1/X nanoseconds cycle time

# How many clock cycles?

---

Number of CPU cycles = Instructions executed \*  
Average Clock Cycles per Instruction (CPI)

*or*

$$\text{CPI} = \text{CPU clock cycles} / \text{Instruction count}$$



# All Together Now

---

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

# All Together Now

---

$$\begin{array}{ccccccc} & \nearrow \text{seconds} & & & \nearrow \text{cycles/instruction} & & \\ \text{CPU Execution} & & \text{Instruction} & \times & \text{CPI} & \times & \text{Clock Cycle} \\ \text{Time} & = & \text{Count} & & & & \text{Time} \\ & \searrow \text{instructions} & & & & & \searrow \text{seconds/cycle} \end{array}$$

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

IC = 40 billion, 4 GHz processor,  
execution time of 30 seconds.

What is the CPI for this program?

Selection	CPI
A	3
B	30
C	1.5
D	$15 \times 10^9$
E	None of the above

# Examples

---

- 4 GHz processor, program runs in 30 seconds, executing 40 billion instructions with CPI above
- New compiler reduces IC to 32 billion, but changes CPI to 3.3: good or bad?

# Putting it together

- Suppose you have a 2 GHz Core i7 called Machine A.
- You also have a 4 GHz Core i7 called Machine B.
- Let's say they have the same underlying architecture (same pipelines, caches, etc.), just one is running at a faster clock rate.
- Running program X (same binary, etc.), could Machine A have an average CPI of 0.9 for program X and Machine B have an average CPI of 1.1 for program X?

Selection	Answer
A	Yes, this is possible but unlikely
B	Yes, this is possible and likely
C	No, this is impossible

# Who Affects Performance?

---

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

- programmer
- compiler
- instruction-set architect
- machine architect
- hardware designer

# What Affects Performance?

---

$$\text{CPU Execution Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

- pipelining
- superpipelining
- cache
- from CISC to RISC
- superscalar

# Key Points

---

- We need to be precise about how to specify performance.
- Performance is only meaningful in the context of a workload.
- Be careful how you summarize performance.
- Amdahl's law
- $ET = IC * CPI * CT$