

Principles of Computer Architecture

CSE 240A

Fall 2024

Hadi Esmaeilzadeh

hadi@ucsd.edu

University of California, San Diego





Modern Times

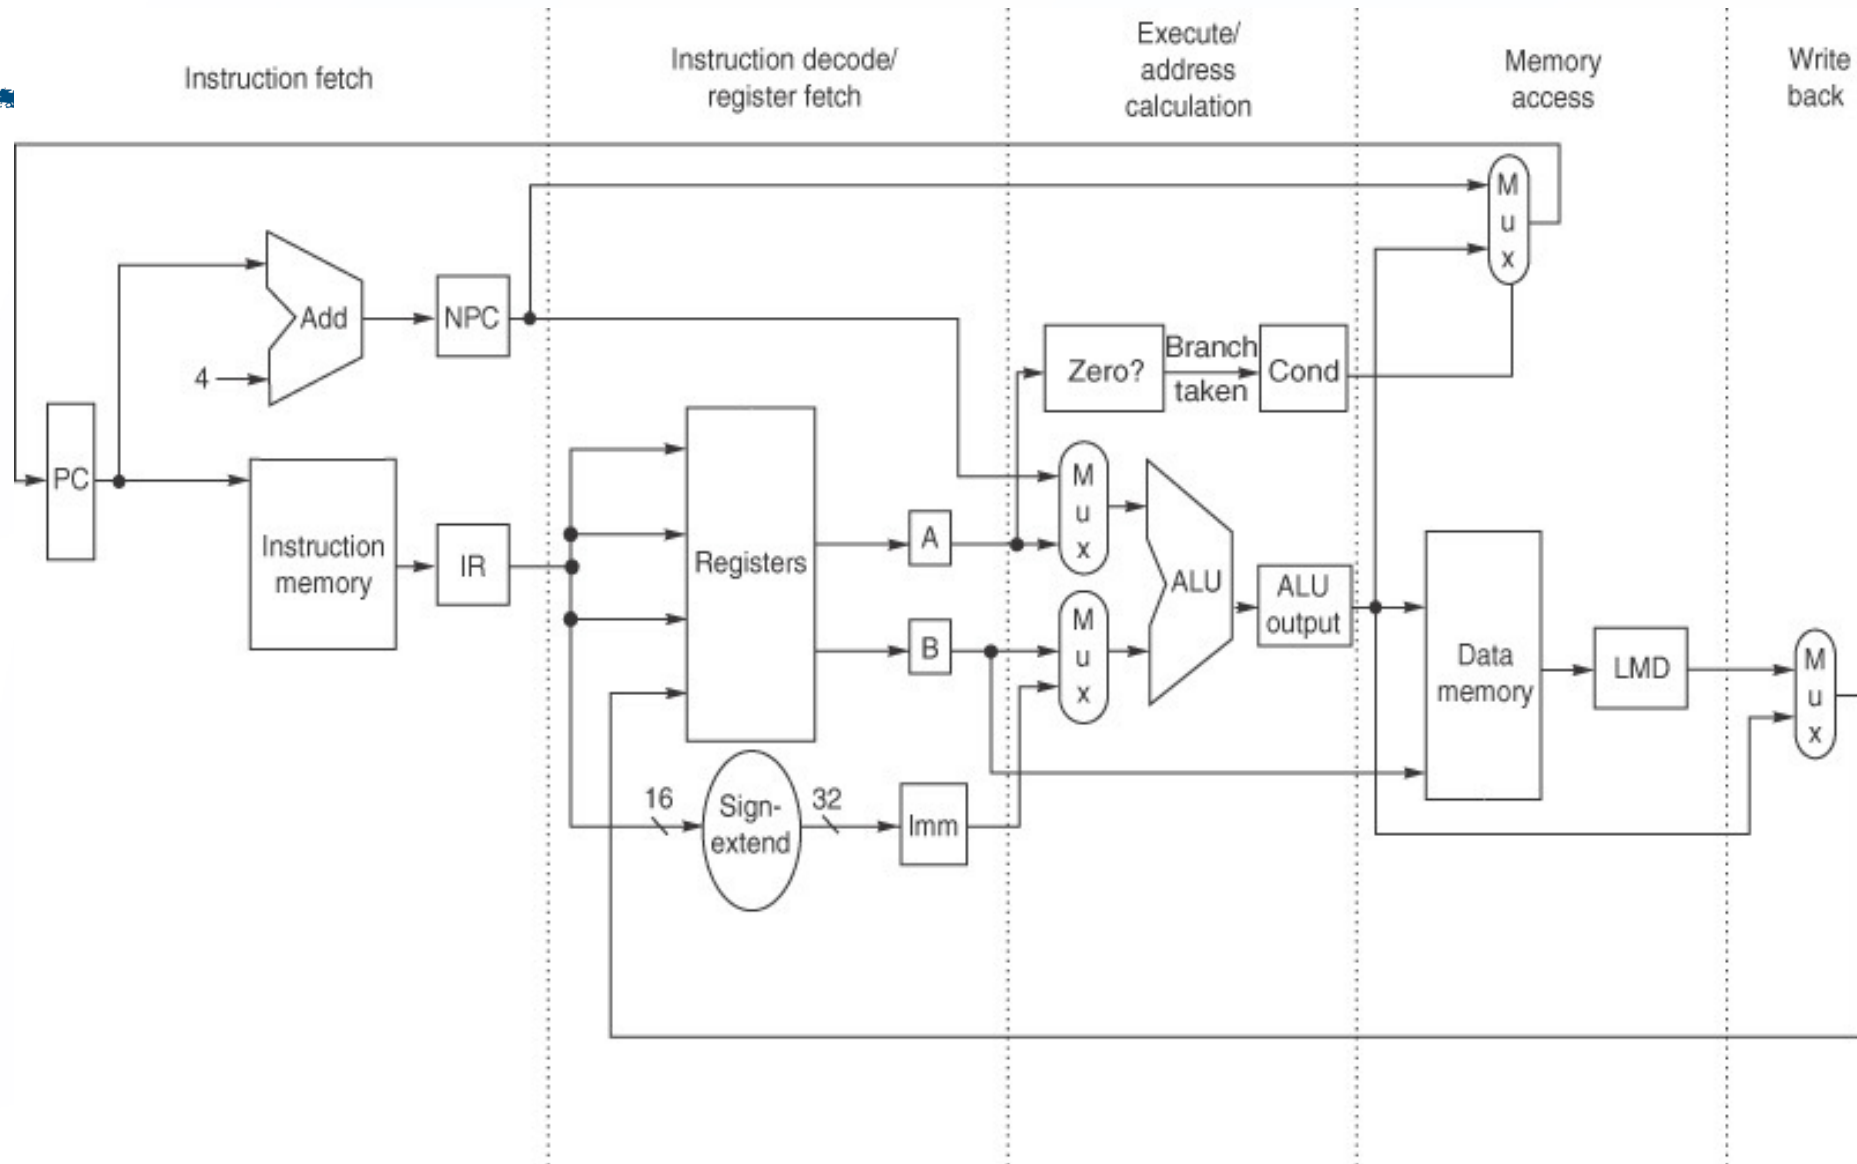
Charlie Chaplin

Modern Times © Roy Export SAS

Pipelining (some general principles)

- Requires separable jobs/stages
- Requires separate resources
- Achieves parallelism without replication
- Pipeline efficiency (keeping the pipeline full) critical to performance

5 Steps of the MIPS Datapath



Notice only
supports
beqz

5 Steps of the MIPS Datapath

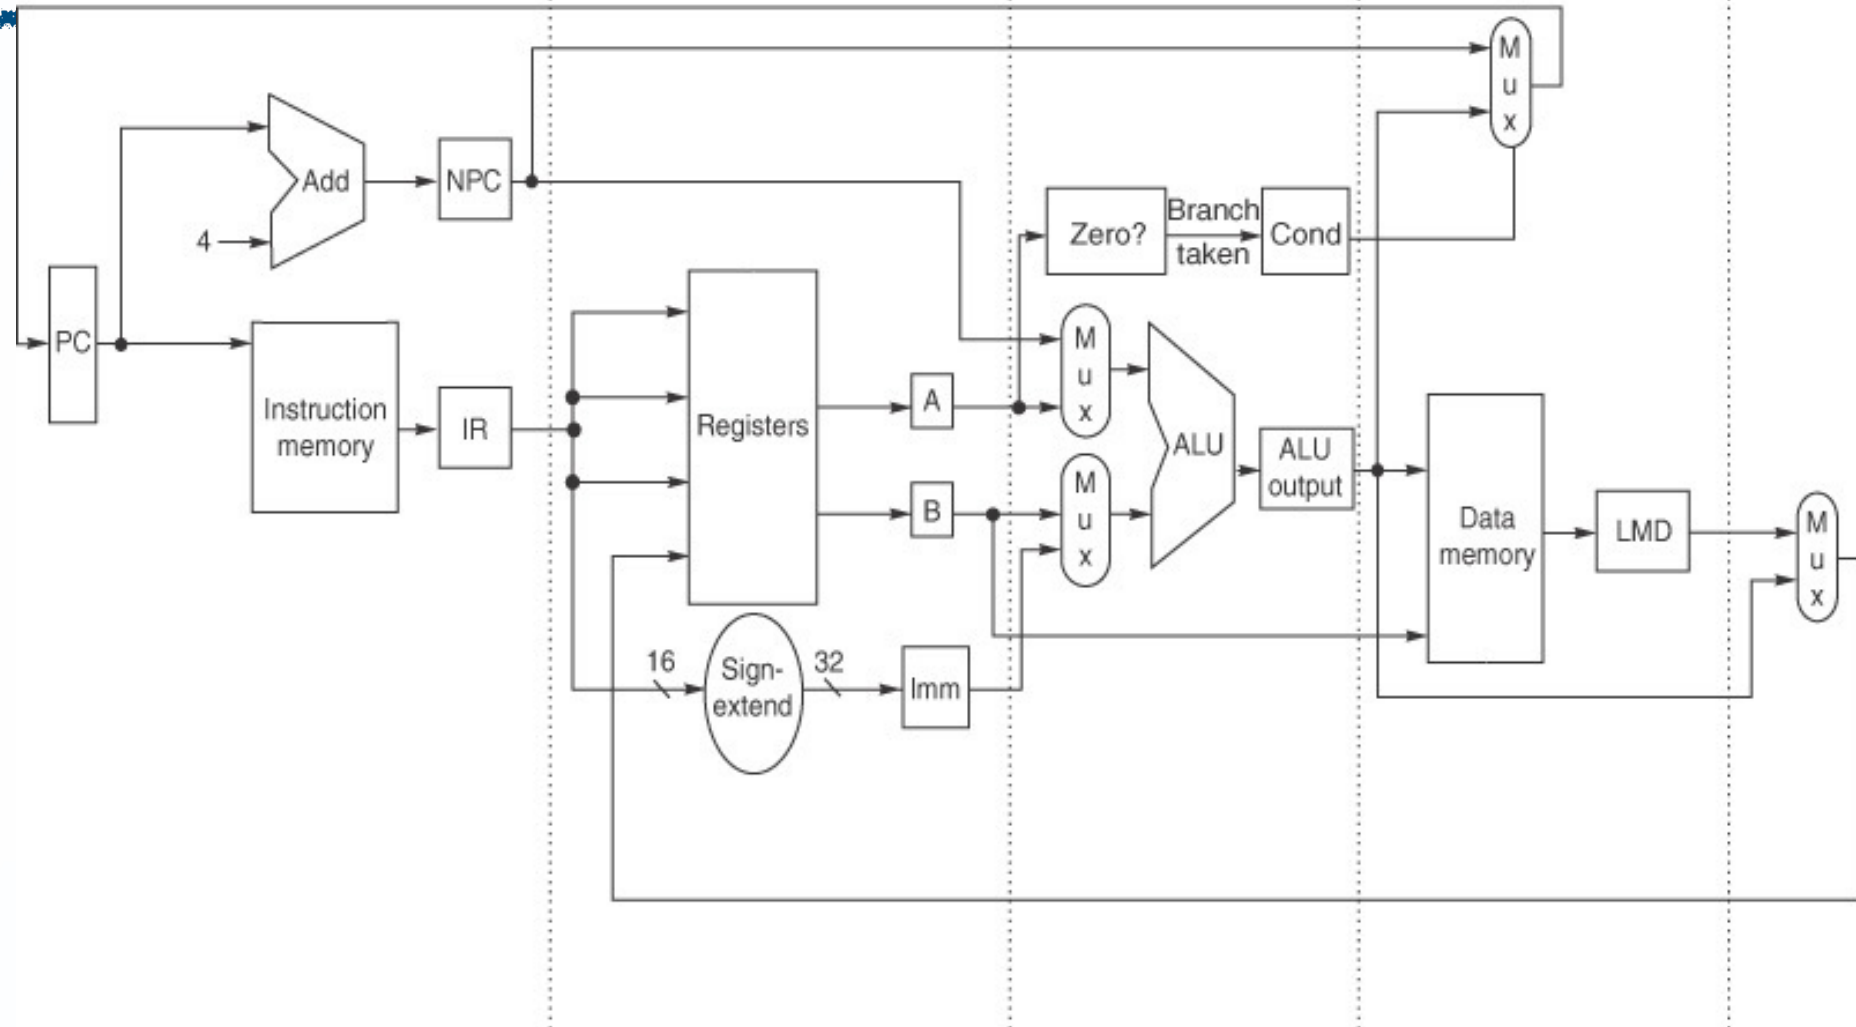
Instruction fetch

Instruction decode/
register fetch

Execute/
address
calculation

Memory
access

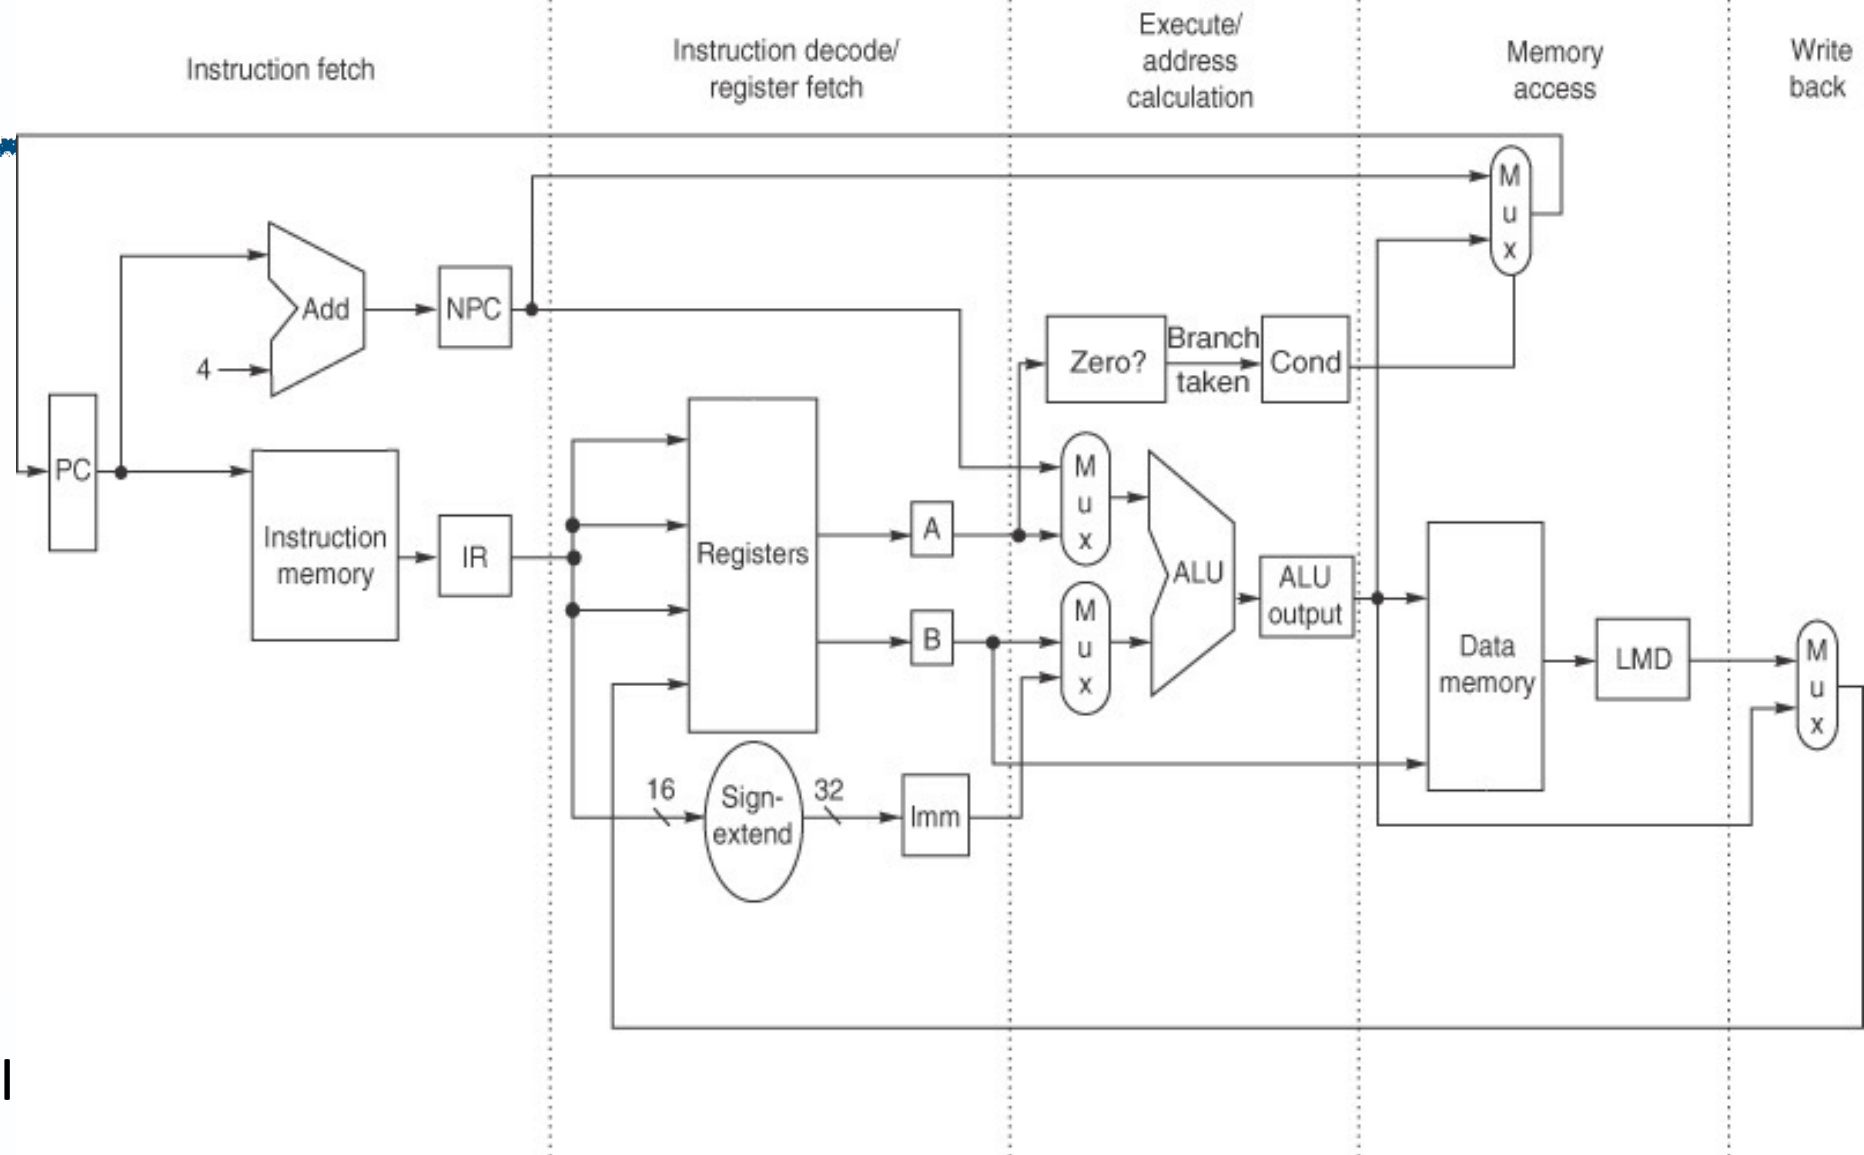
Write
back



Instruction Fetch (IF)

- $IR \leftarrow M[PC]$
- $NPC \leftarrow PC + 4$

5 Steps of the MIPS Datapath



Instruction Decode/ register read (**ID**)

- $A \leftarrow \text{Reg}[\text{IR6}..10]$
- $B \leftarrow \text{Reg}[\text{IR11}..15]$
- $\text{Imm} \leftarrow \text{sign_extend}(\text{I}$

5 Steps of the MIPS Datapath

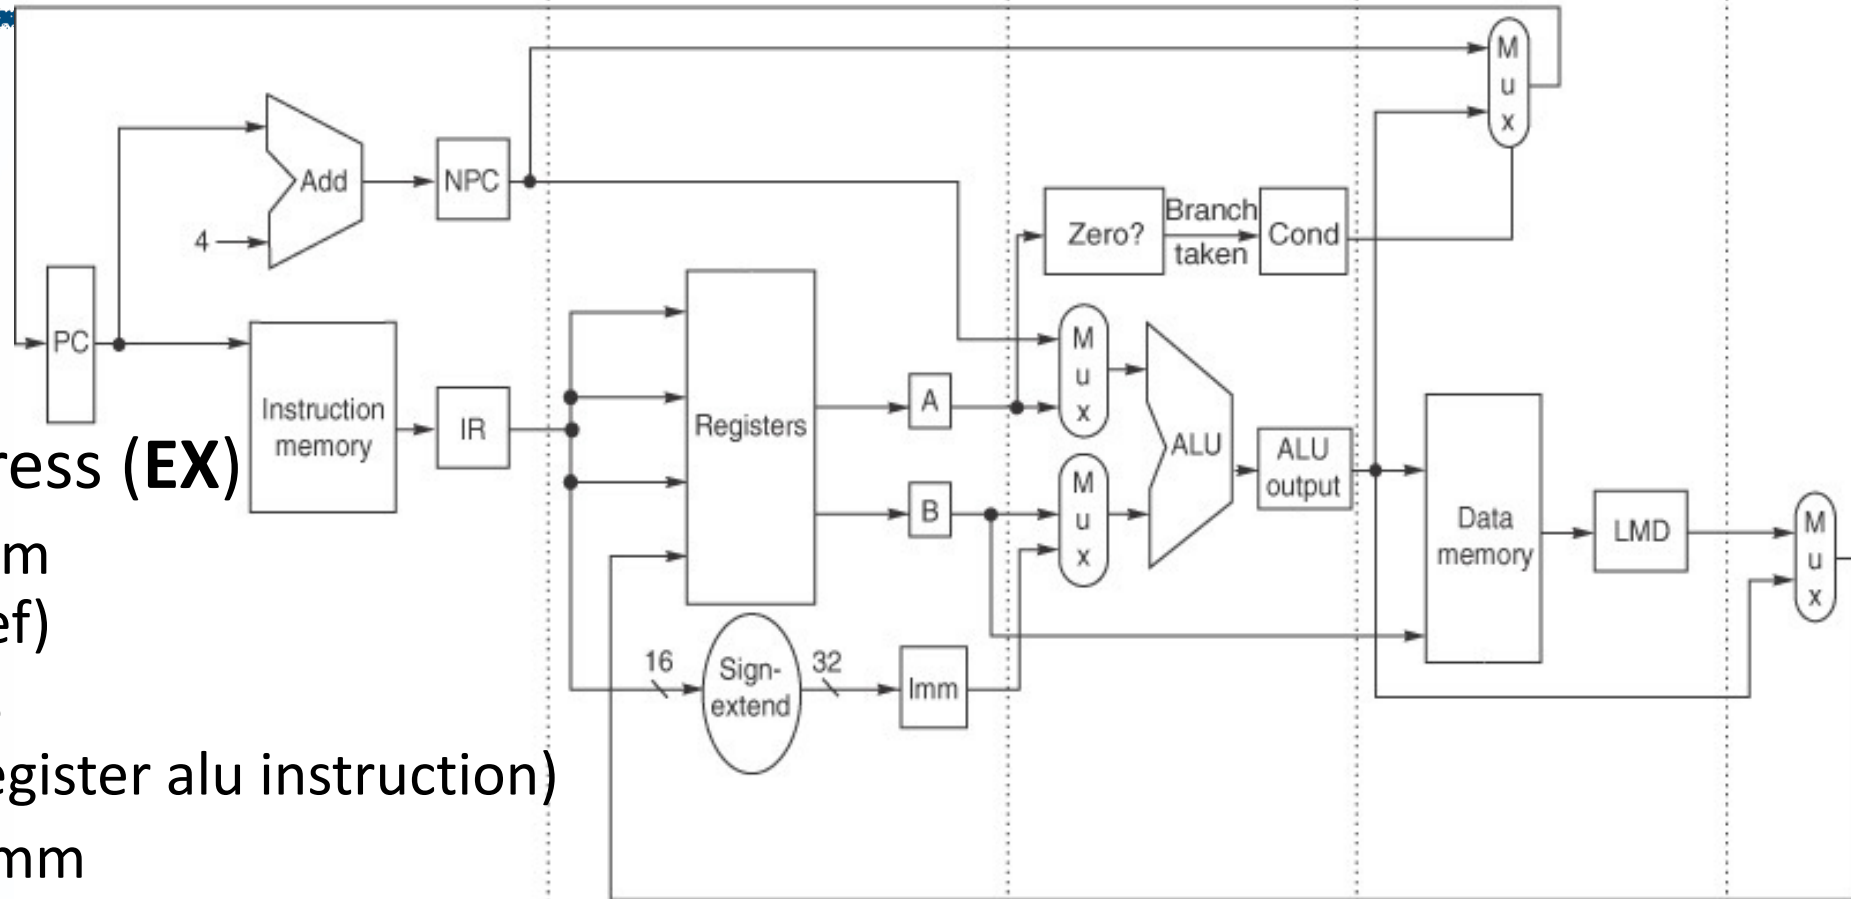
Instruction fetch

Instruction decode/
register fetch

Execute/
address
calculation

Memory
access

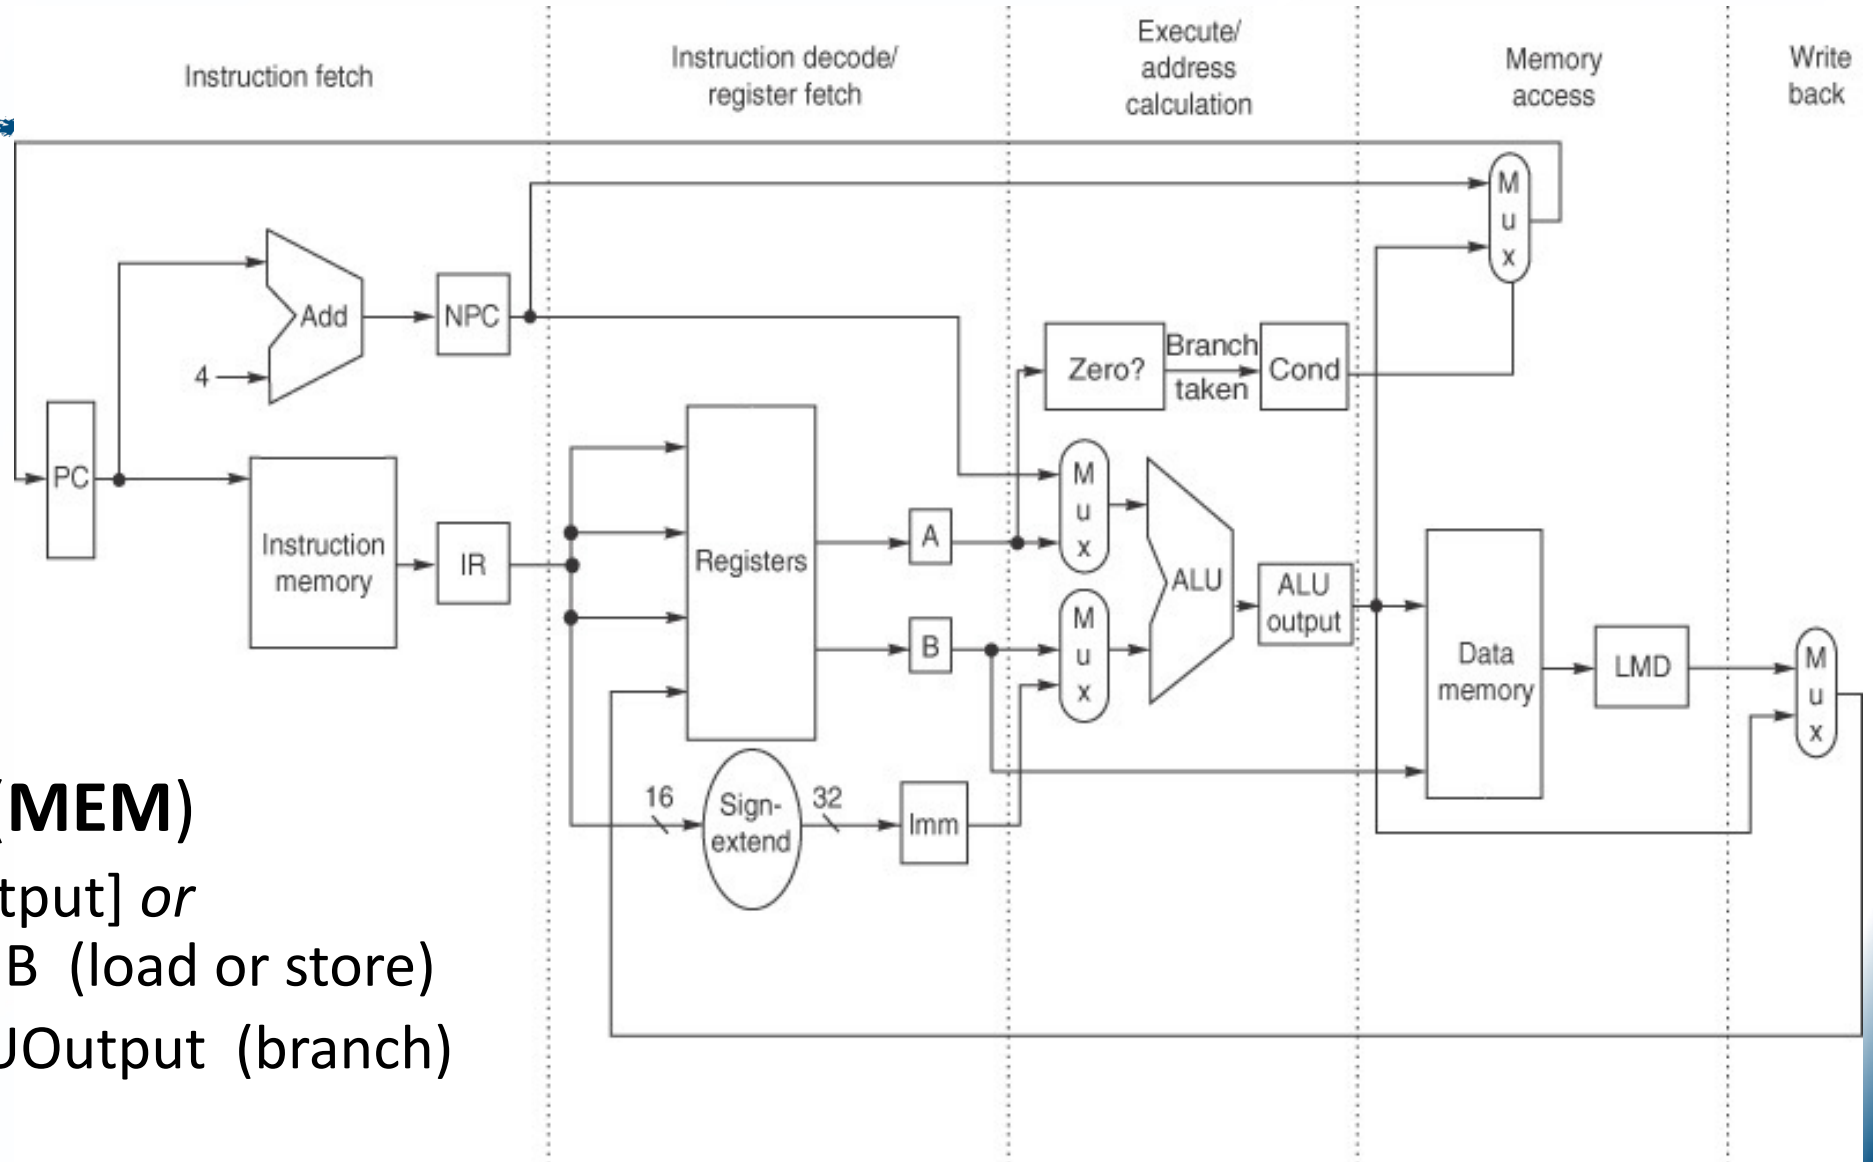
Write
back



Execute/Effective Address (EX)

- $\text{ALUOutput} \leftarrow A + \text{Imm}$
(memory ref)
- $\text{ALUOutput} \leftarrow A \text{ op } B$
(register-register alu instruction)
- $\text{ALUOutput} \leftarrow A \text{ op Imm}$
(register-immediate alu instruction)
- $\text{ALUOutput} \leftarrow \text{NPC} + \text{Imm}; \text{Cond} \leftarrow (A \text{ op } 0)$
(Branch)

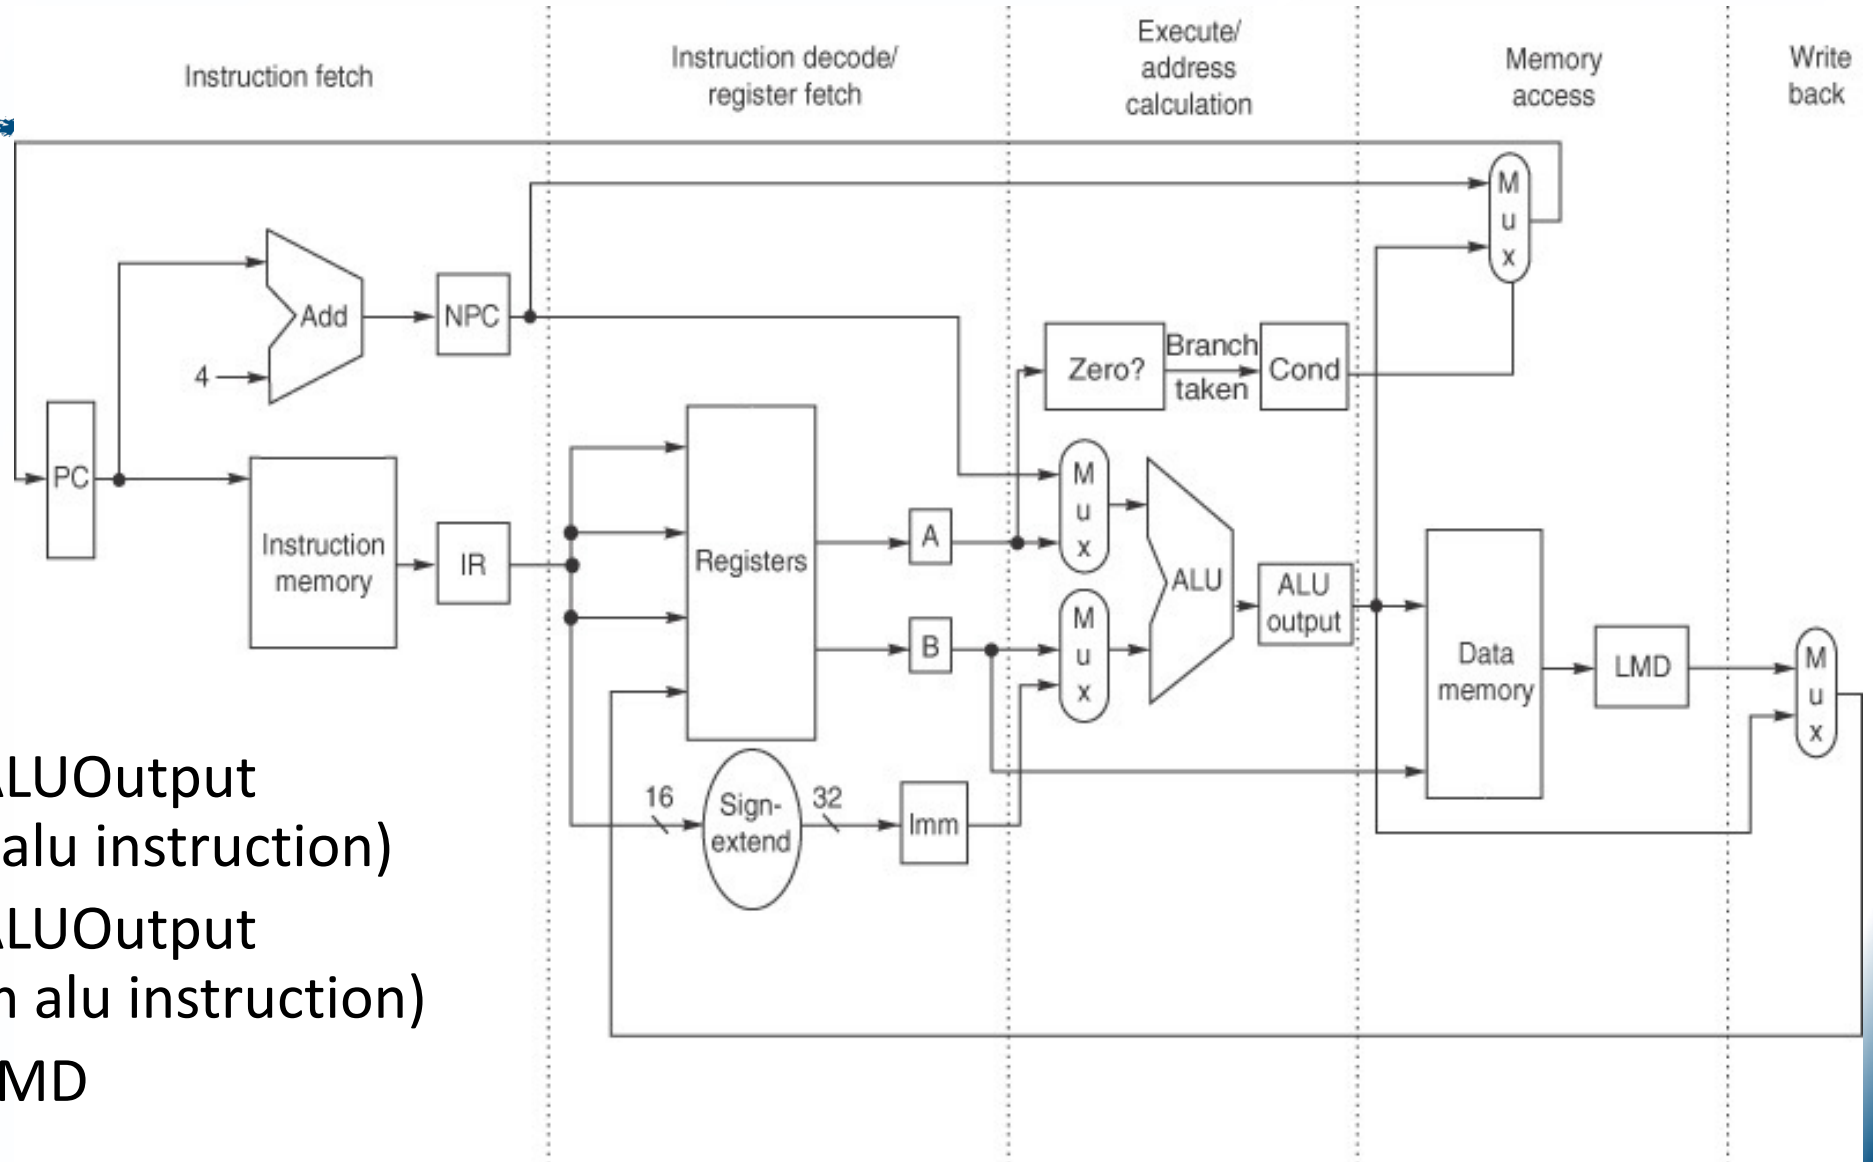
5 Steps of the MIPS Datapath



Memory access/ branch completion (**MEM**)

- $LMD \leftarrow M[ALUOutput]$ or $M[ALUOutput] \leftarrow B$ (load or store)
- if (cond) $PC \leftarrow ALUOutput$ (branch)
else $PC \leftarrow NPC$

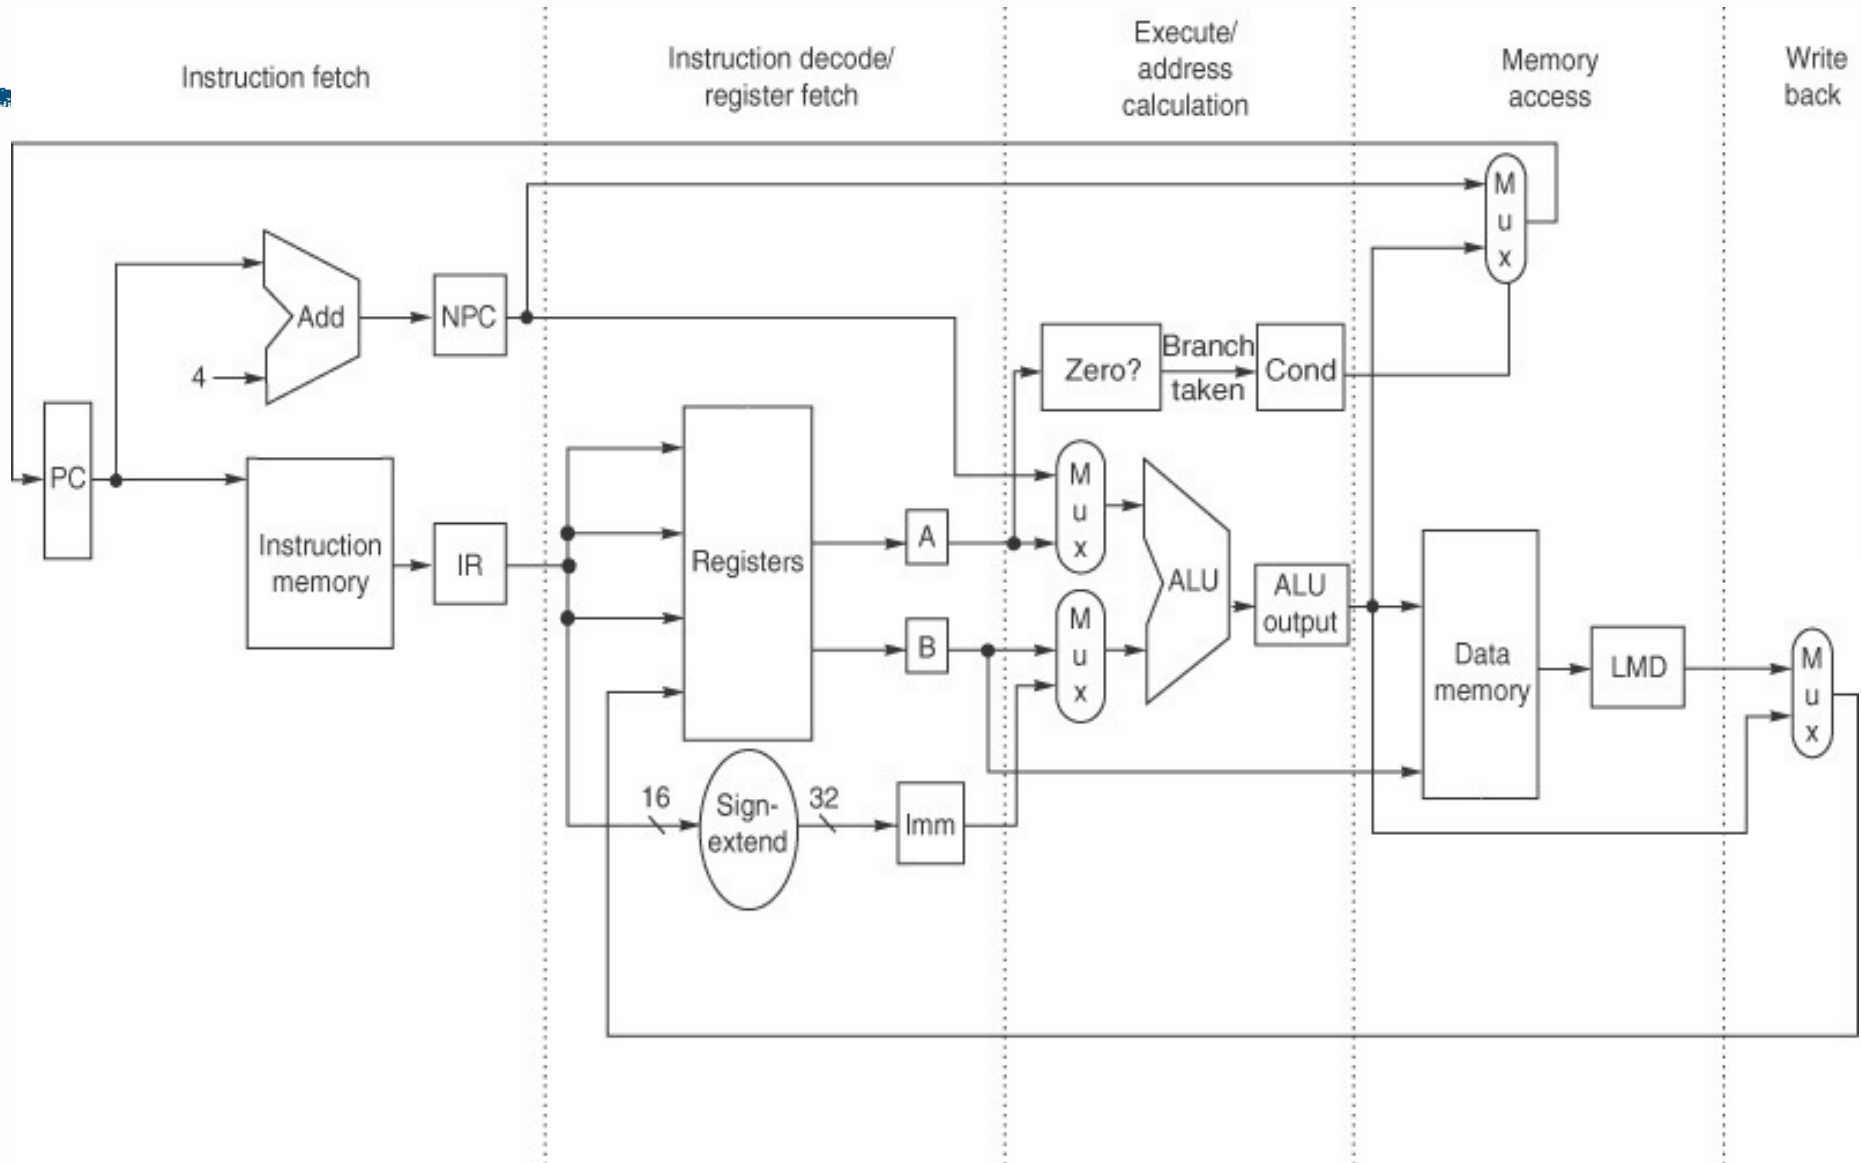
5 Steps of the MIPS Datapath




Write-Back (WB)

- $\text{Reg}[\text{IR16}..\text{20}] \leftarrow \text{ALUOutput}$
(reg-reg alu instruction)
- $\text{Reg}[\text{IR11}..\text{15}] \leftarrow \text{ALUOutput}$
(reg-imm alu instruction)
- $\text{Reg}[\text{IR11}..\text{15}] \leftarrow \text{LMD}$

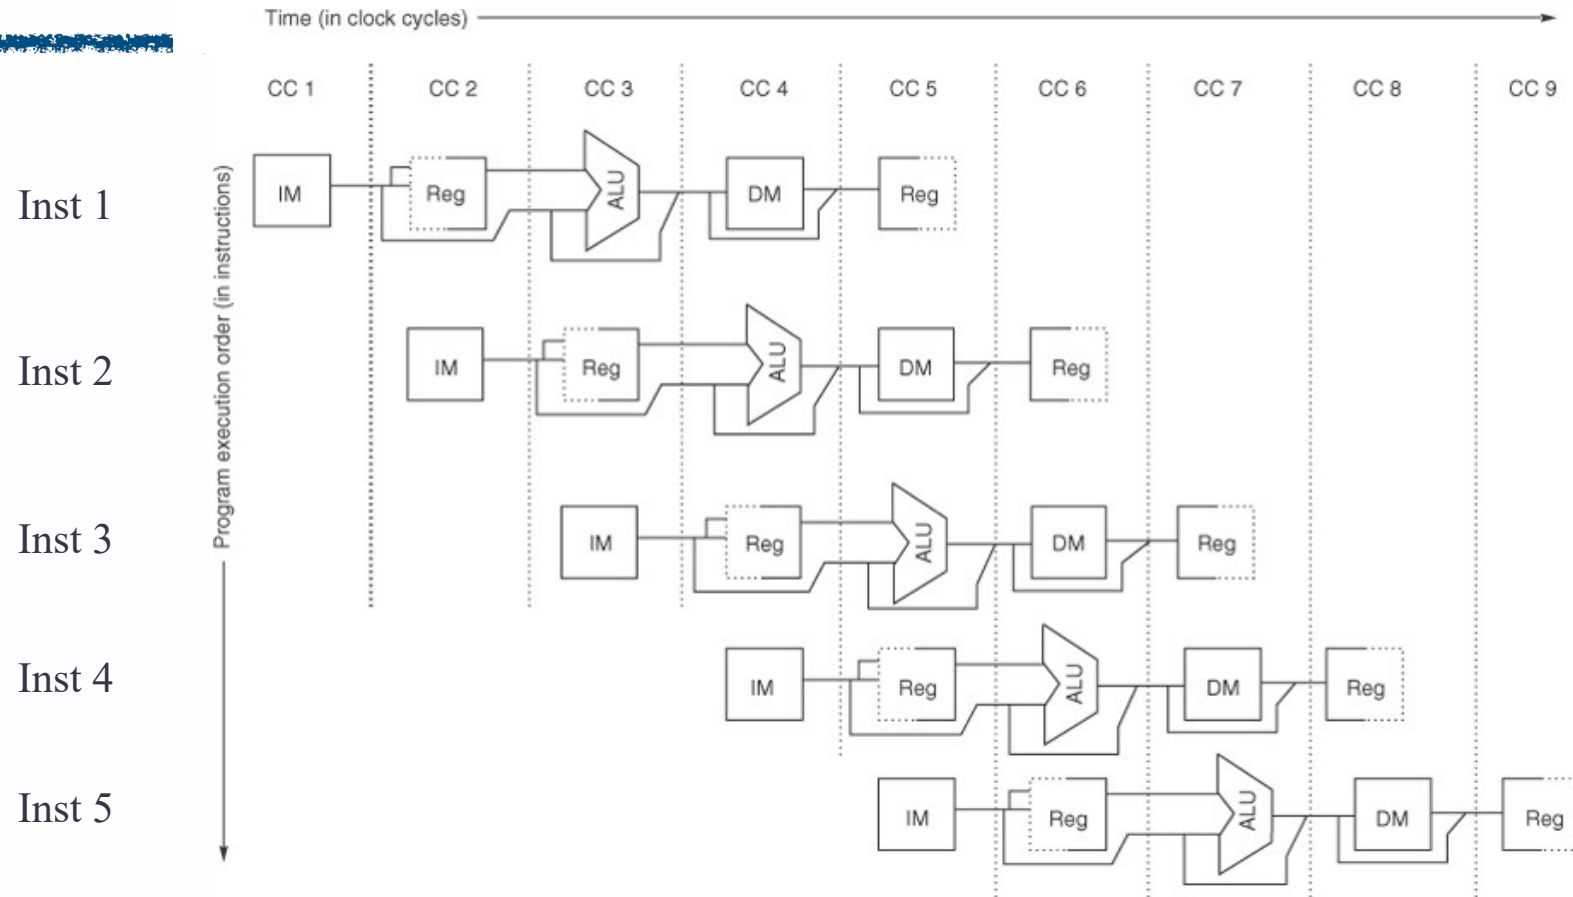
ADDI R7, R2, #35



5 Steps of a MIPS Instruction



Visualizing Pipelining

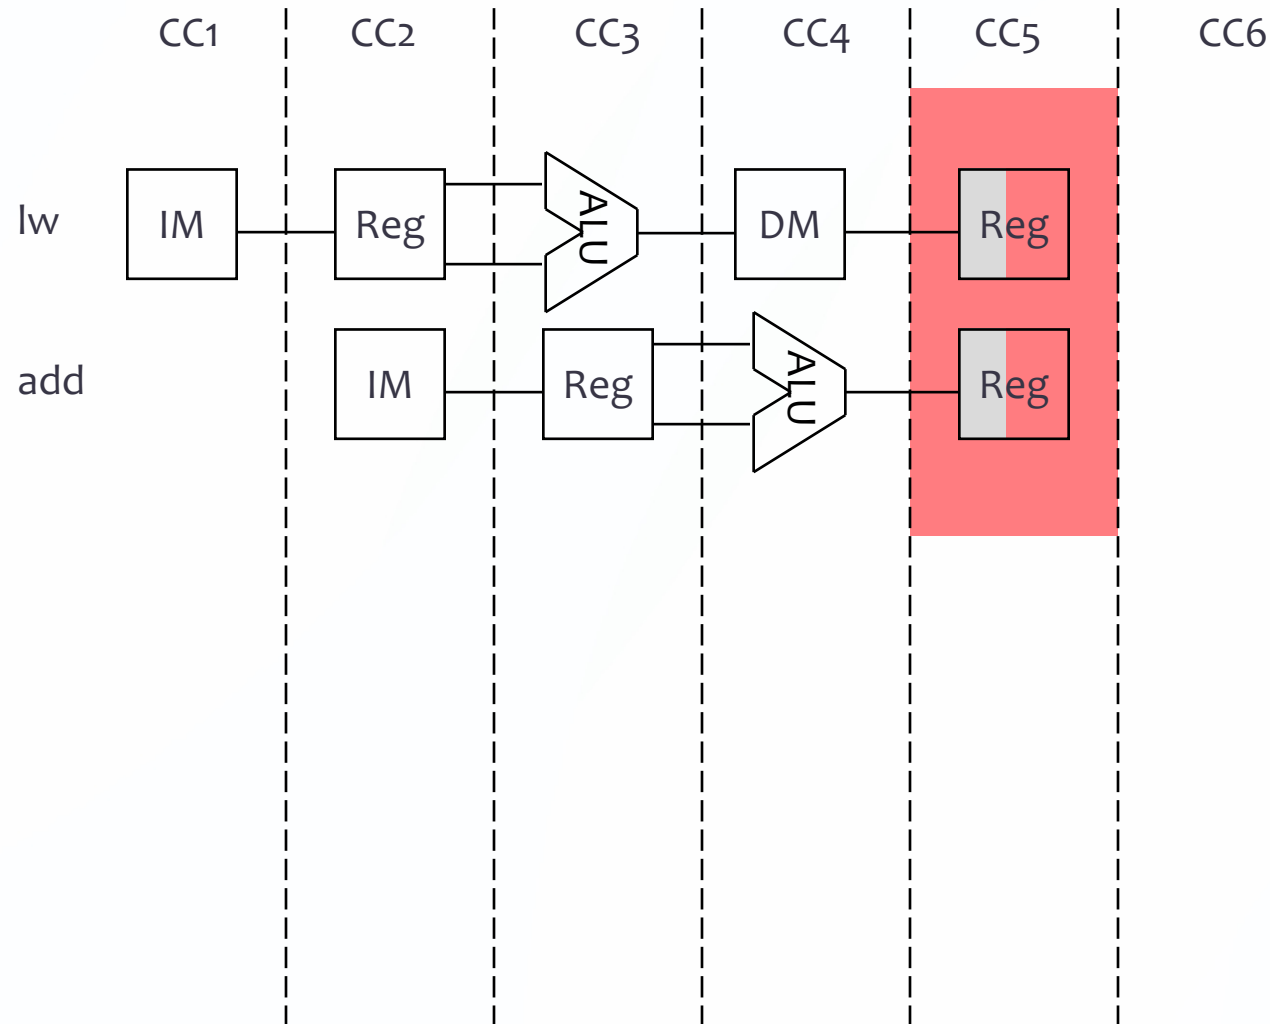


Pipeline Stages

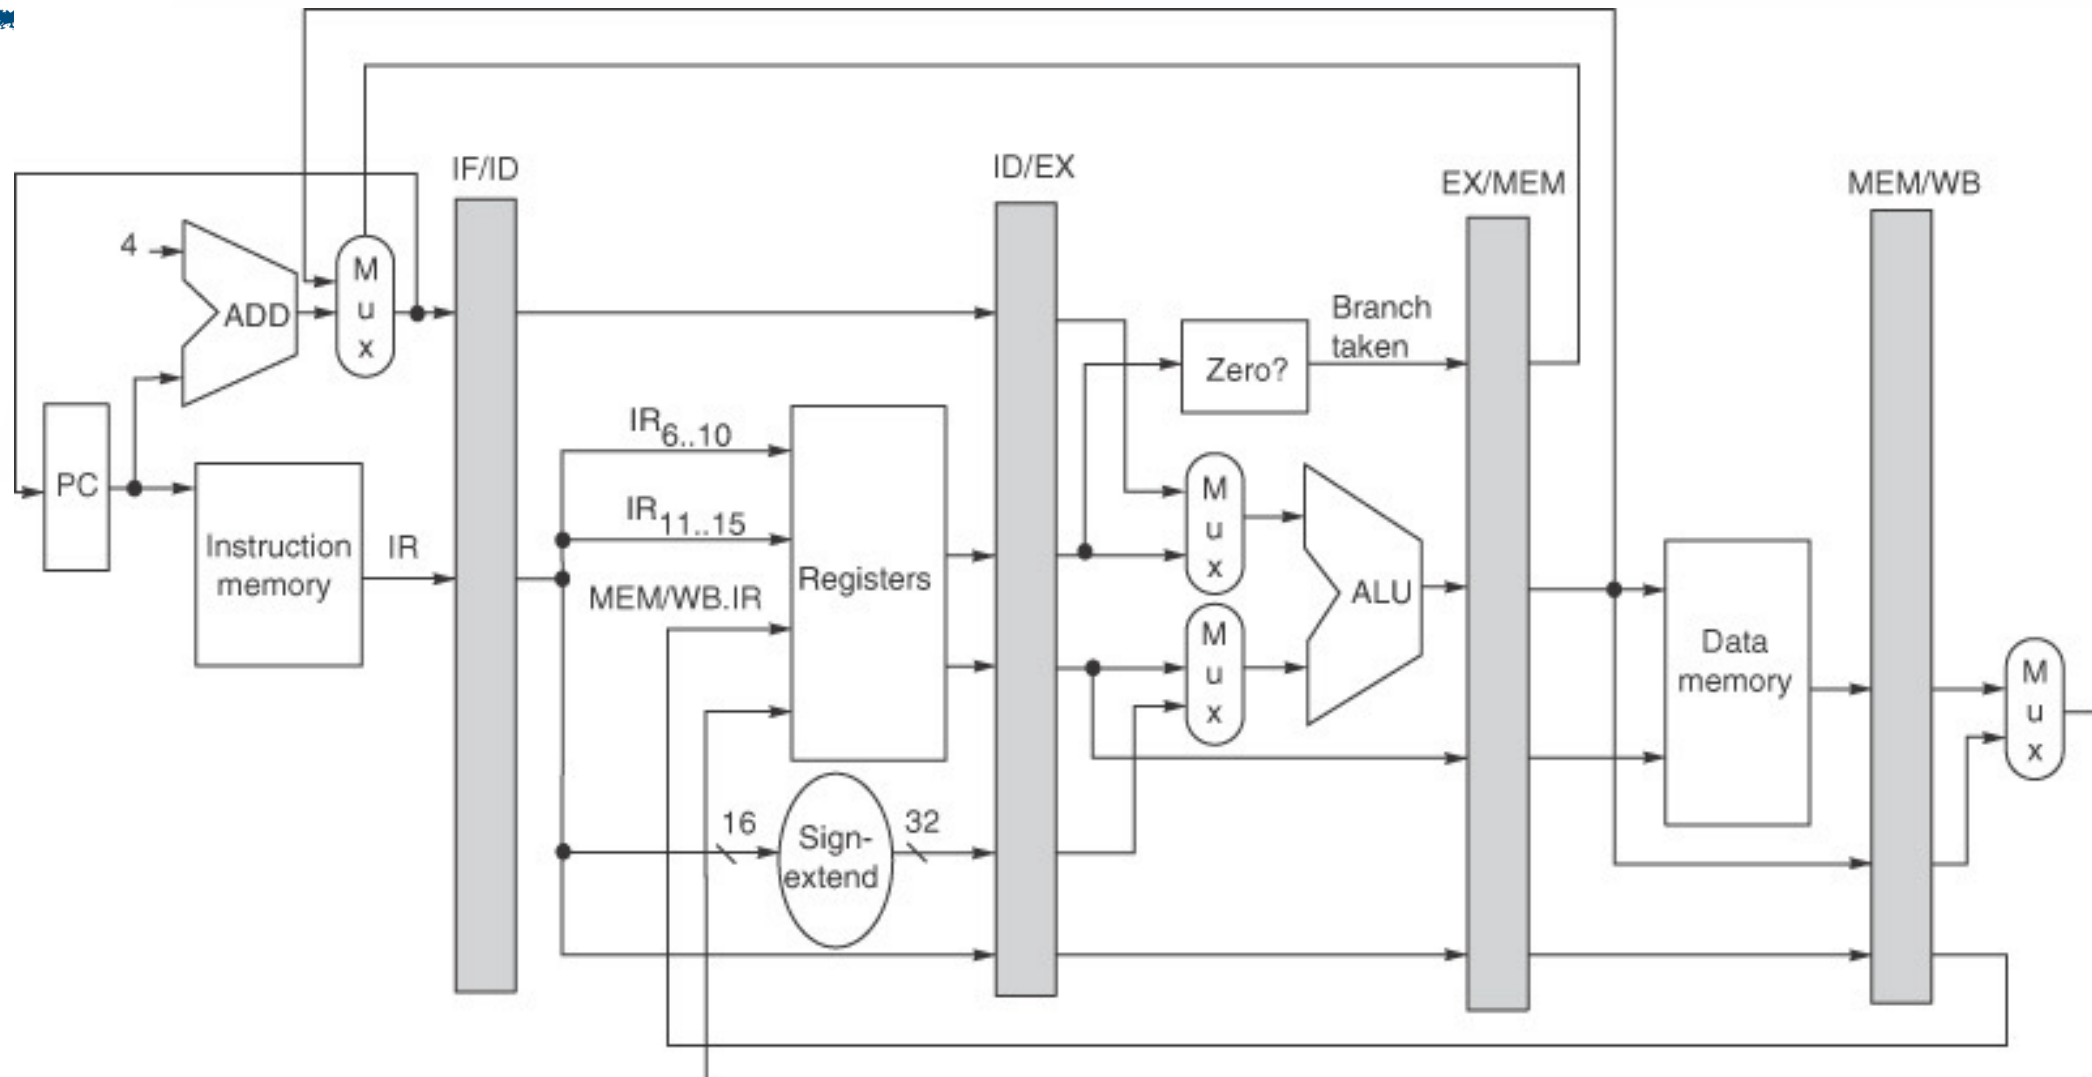
Instead of forcing every instruction to go through all 5 stages, could we let them finish early if done early? For example, an R-type instruction can finish in 4 cycles, can we just let it finish early?

Selection	Yes/No	Reason (Choose BEST answer)
A	Yes	Decreasing R-type to 4 cycles improves instruction throughput
B	Yes	Decreasing R-type to 4 cycles improves instruction latency
C	No	Decreasing R-type to 4 cycles causes hazards
D	No	Decreasing R-type to 4 cycles causes hazards and doesn't impact throughput
E	No	Decreasing R-type to 4 cycles causes hazards and doesn't impact latency

Mixed Instructions in the Pipeline



The Pipelined MIPS Datapath

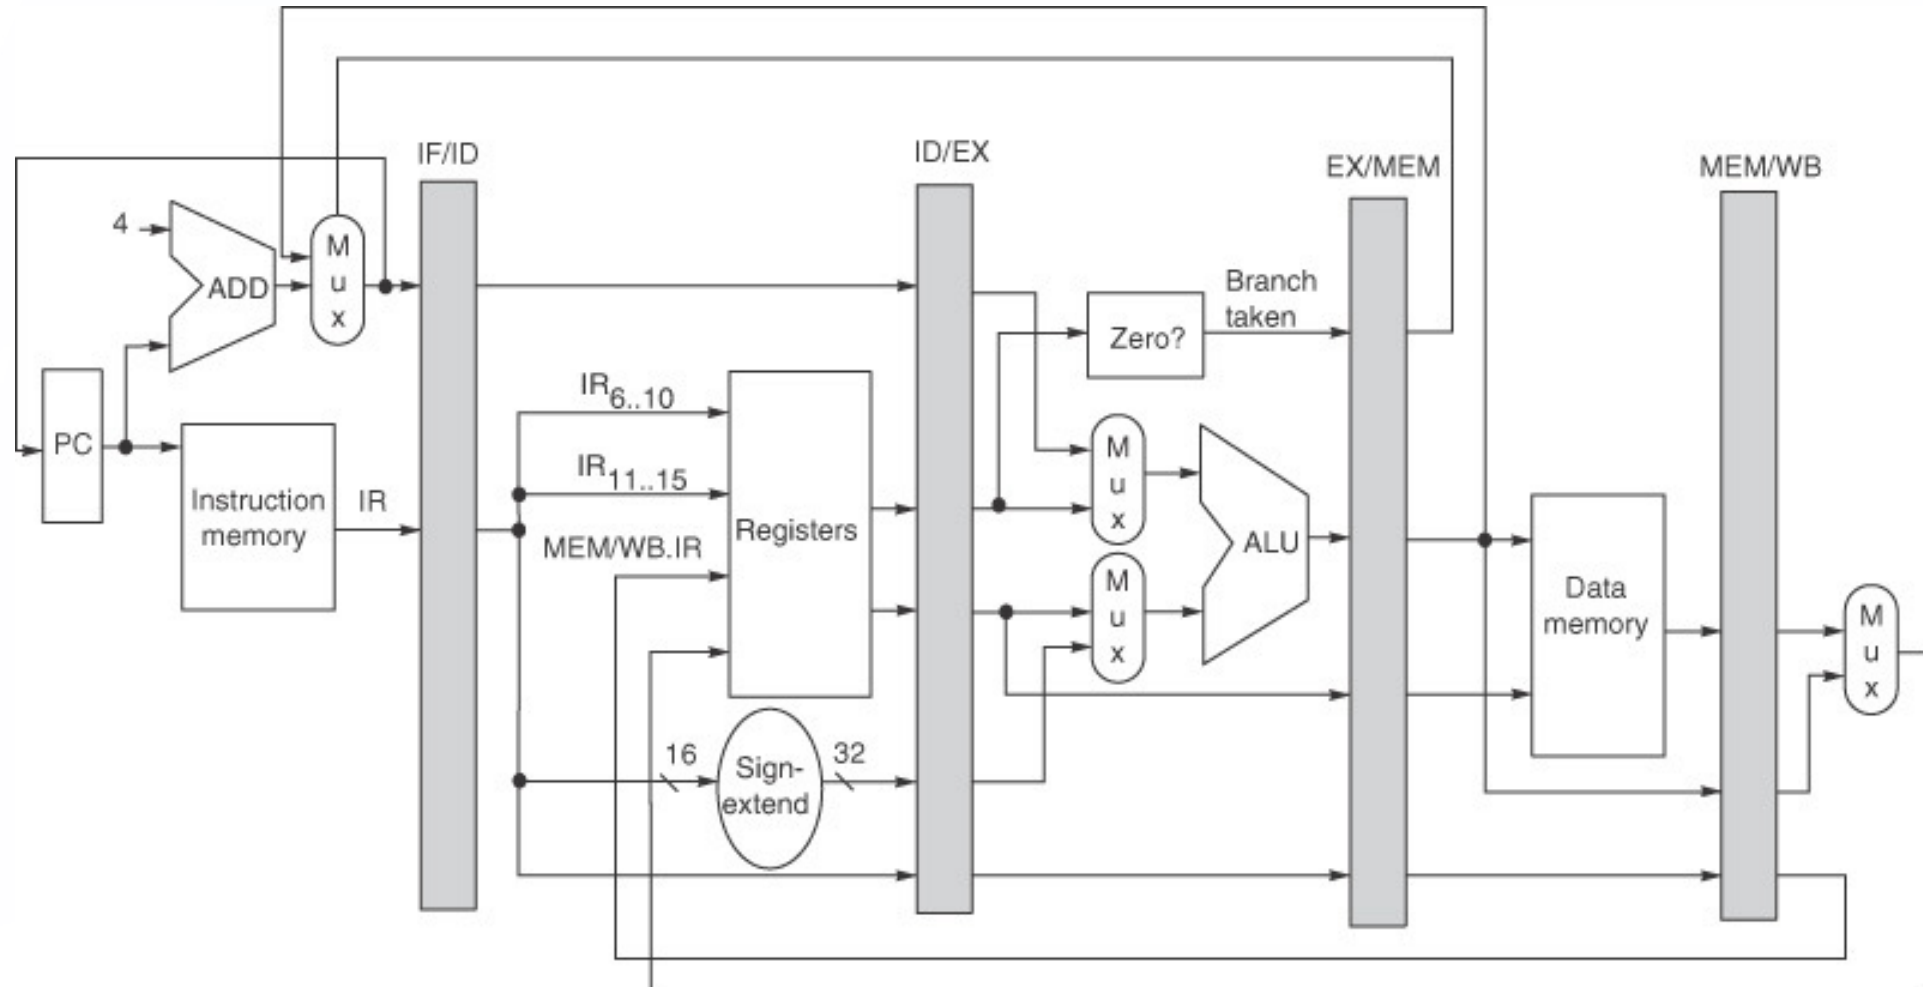


The Pipeline in Motion

- `addi R5, R1, #35`
- `add R6, R2, R1`
- `lw R8, 10000(R3)`

The Pipeline In Motion

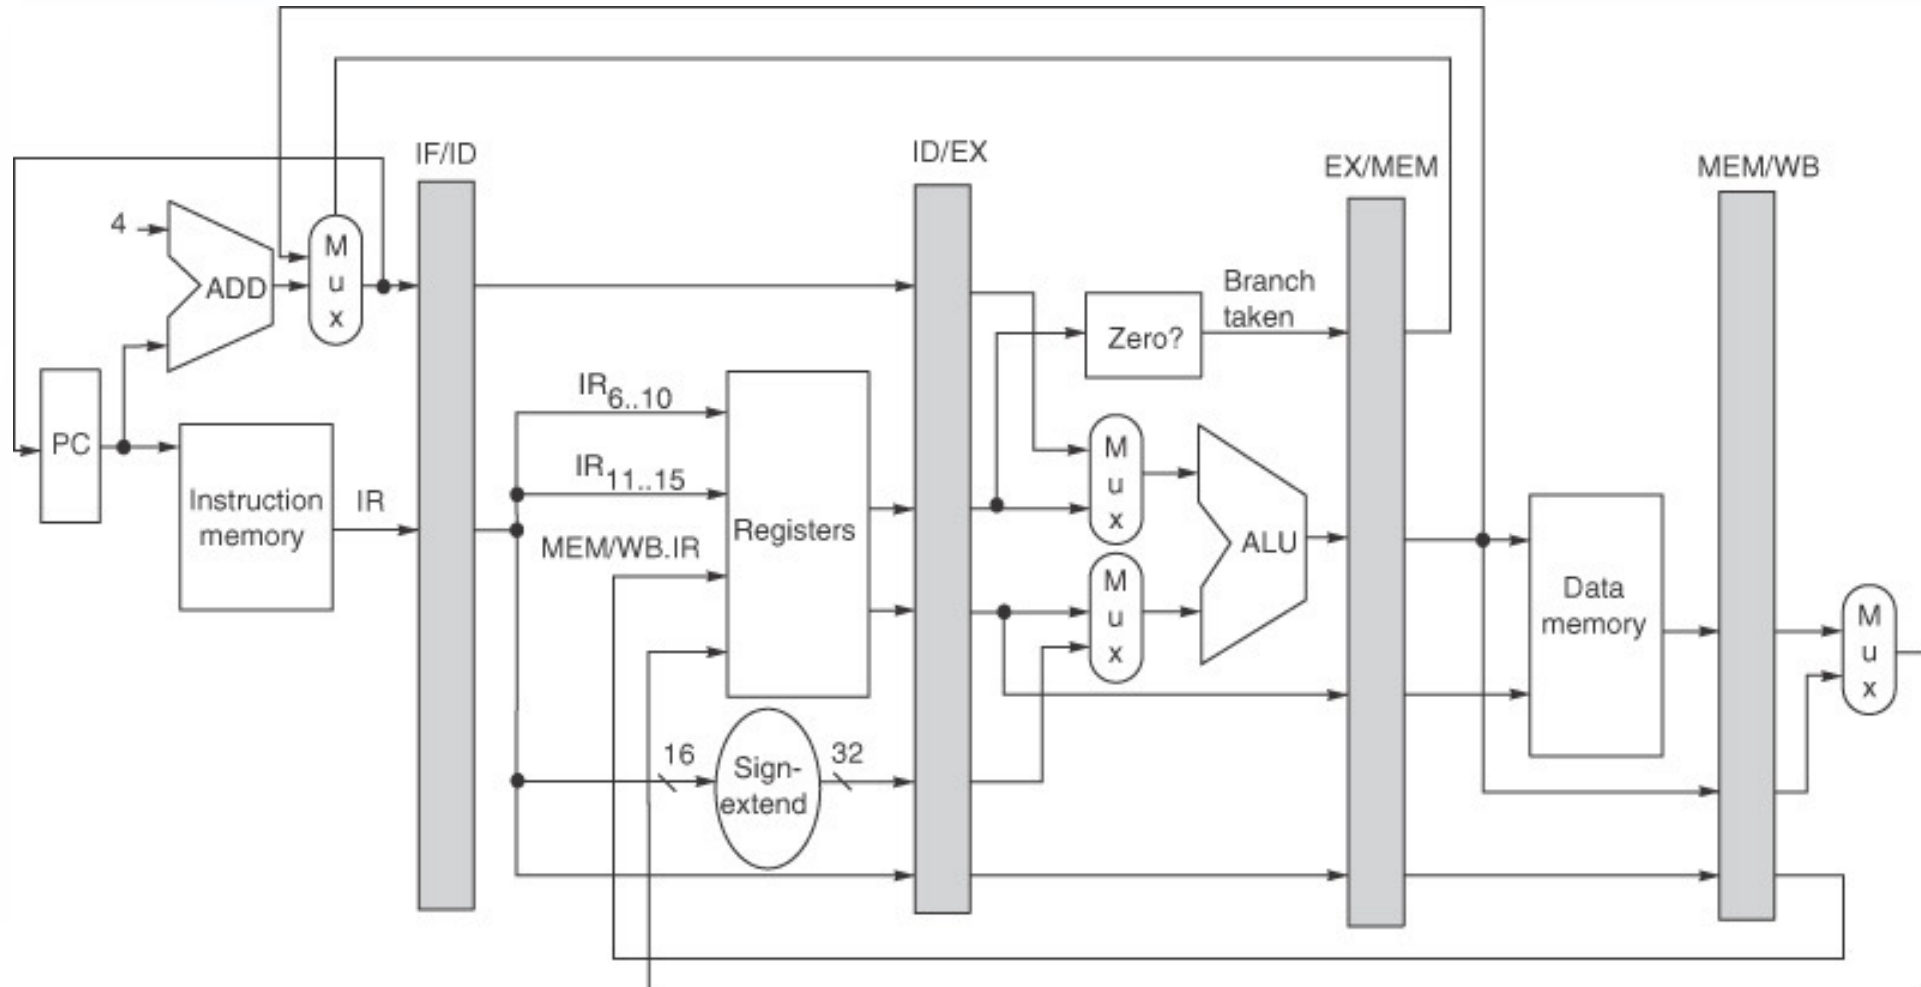
addi R5, R1, #35



The Pipeline In Motion

add R6, R2, R1

addi R5, R1, #35




The Pipeline In Motion

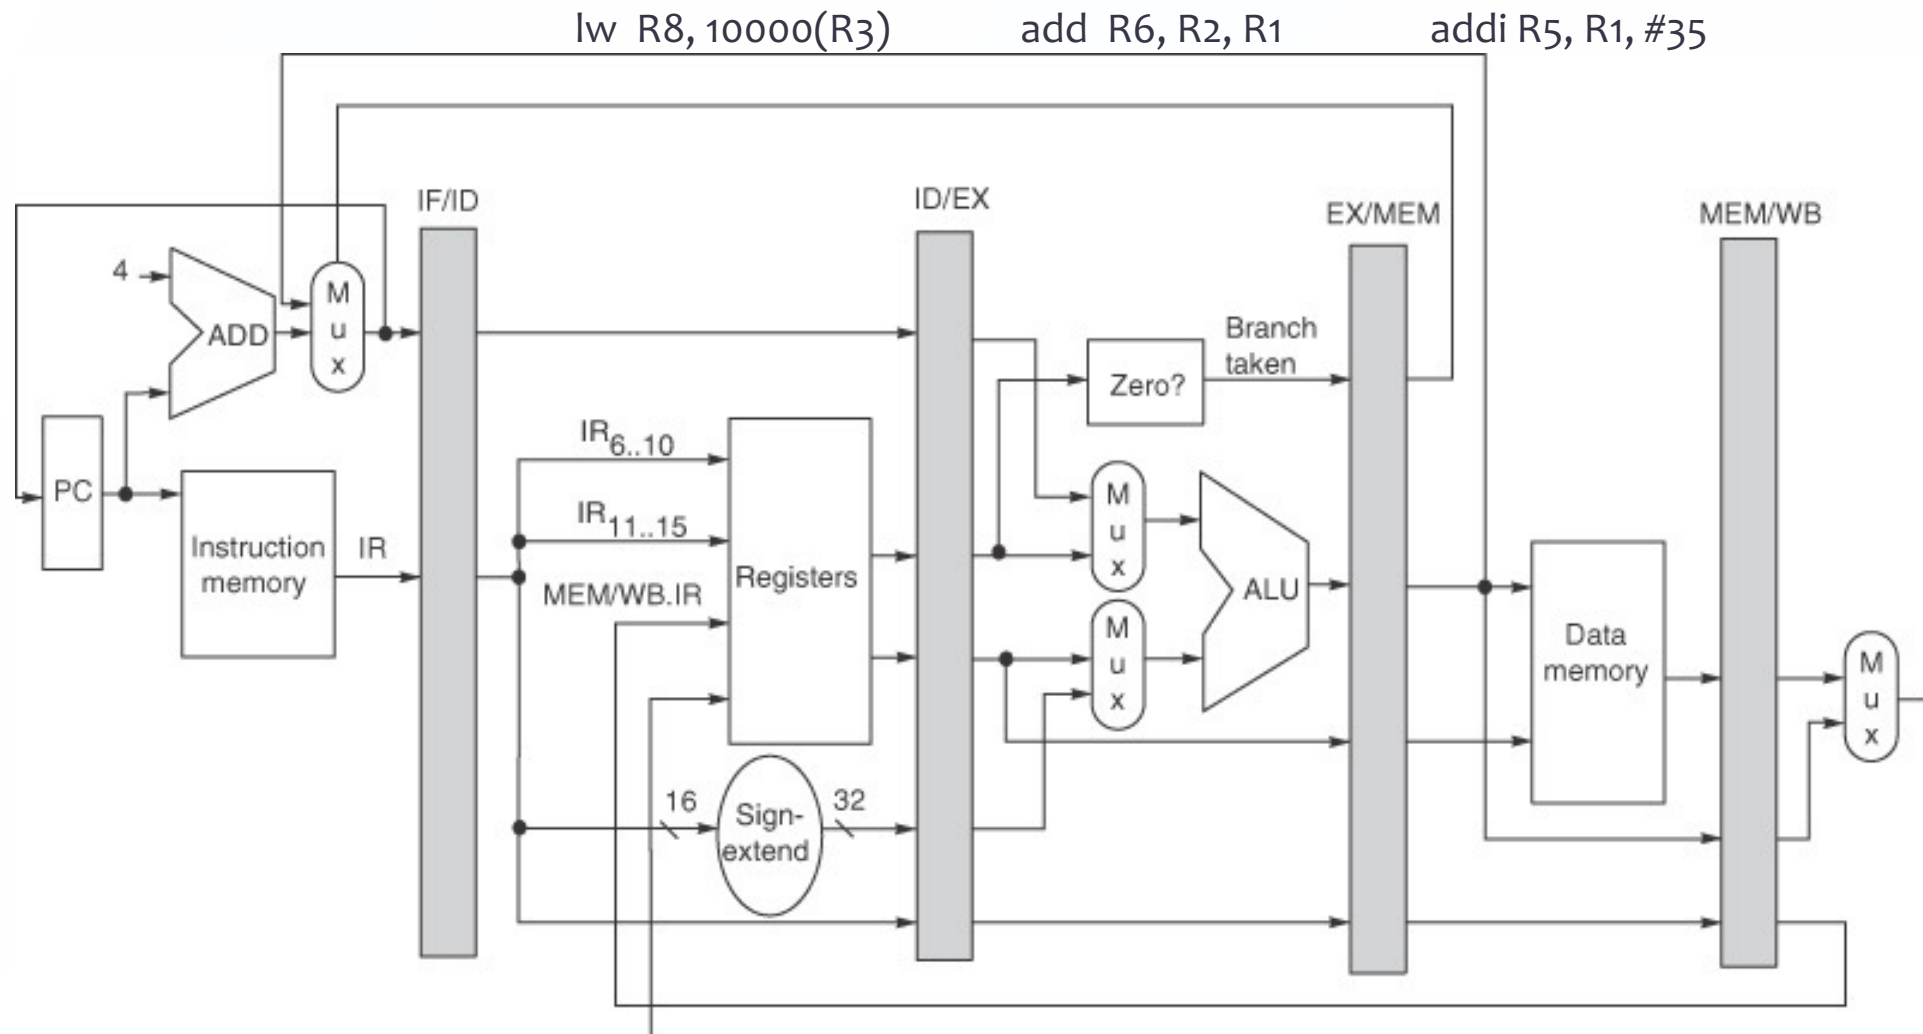
lw R8, 10000(R3)

add R6, R2, R1

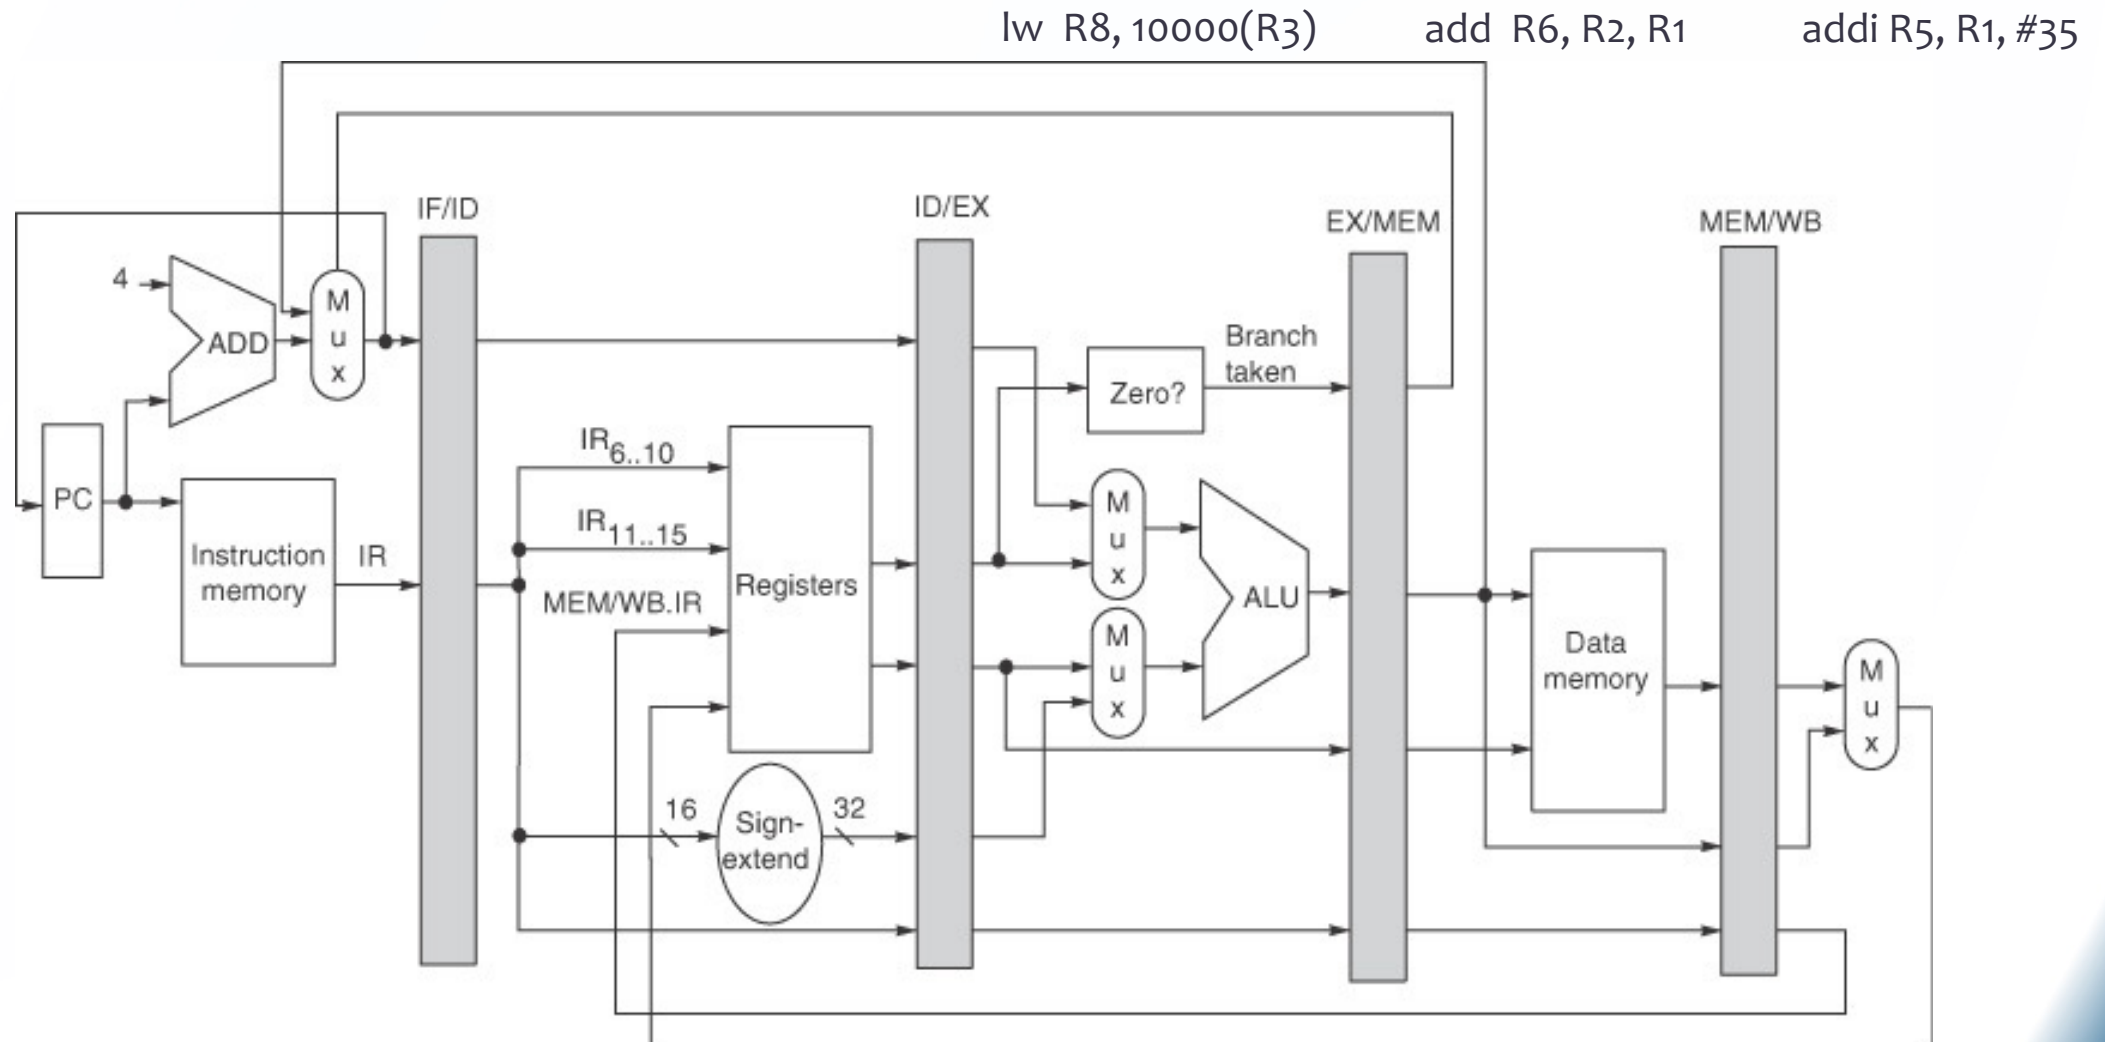
addi R5, R1, #35



The Pipeline In Motion



The Pipeline In Motion



The Pipeline in Motion

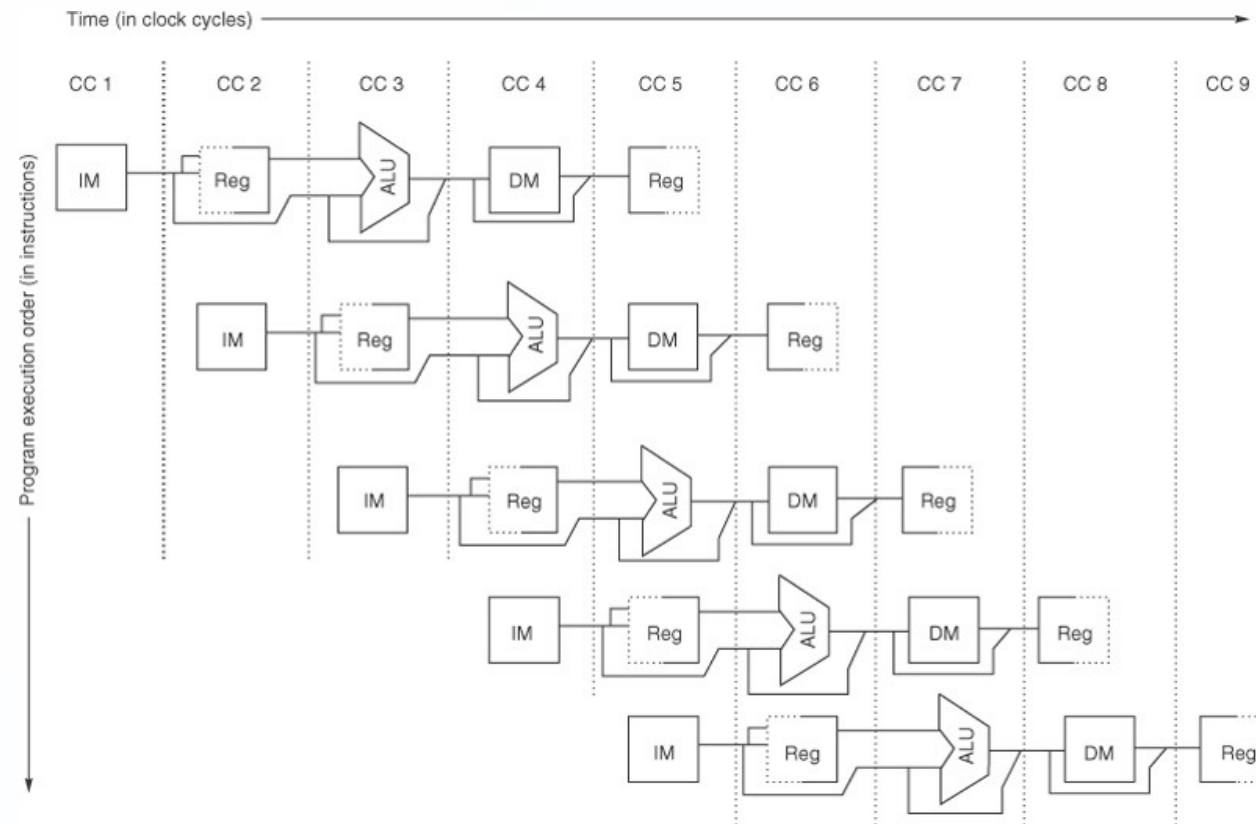
addi R5, R1, #35

add R6, R2, R1

lw R8, 10000(R3)

Inst 4

Inst 5



Instruction Latencies and Throughput

Which of the statements below is true about a pipelined processor?

Selection	Statement
A	Individual instruction latency remains essentially unchanged from single-cycle (excluding overheads); Instruction throughput increases
B	Individual instruction latency remains essentially unchanged from multi-cycle (excluding overheads); Instruction throughput decreases
C	Individual instruction latency improves by a factor of 5 over single-cycle (excluding overheads); Instruction throughput increases
D	Individual instruction latency improves by a factor of 5 over multi-cycle (excluding overheads); Instruction throughput decreases
E	None of the above

Single Cycle vs. Multi-cycle

	CPI	CT
Single Cycle		
Multi-cycle		
Pipeline		

Pipeline Performance

- $ET = IC * CPI * CT$
 - single-cycle processor
 - multiple-cycle processor
 - pipelined processor
- complexity has a cost
 - e.g., latch overhead
 - uneven stage latencies
- Can't always keep the pipeline full
 - why not?

When Things Go Wrong -- Pipeline Hazards

- Limits to pipelining: **Hazards** prevent next instruction from executing during its designated clock cycle
 - *Structural hazards*: HW cannot support this combination of instructions
 - *Data hazards*: Instruction depends on result of prior instruction still in the pipeline
 - *Control hazards*: Pipelining of branches & other instructions that change the PC
- Common solution is to stall the pipeline until the hazard is resolved, inserting one or more “bubbles” in the pipeline

Key Points

Pipeline improves throughput rather than latency

Pipelining gets parallelism without replication

$$ET = IC * CPI * CT$$