

# Principles of Computer Architecture

---

CSE 240A

Fall 2024

Hadi Esmaeilzadeh

[hadi@ucsd.edu](mailto:hadi@ucsd.edu)

University of California, San Diego





















# Processor Cheat Sheet

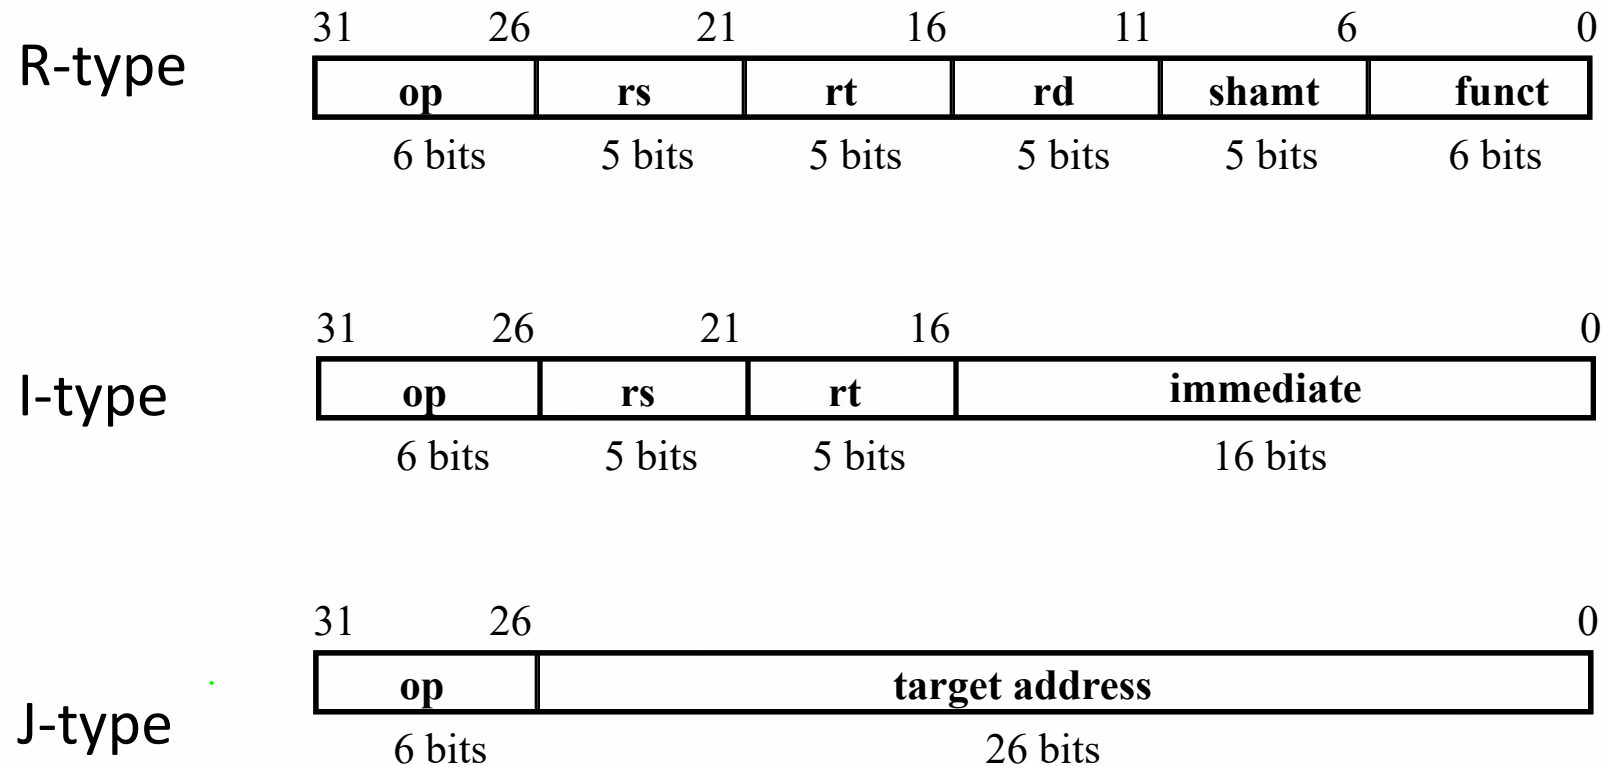
---

# The Processor: Datapath & Control

- We're ready to look at an implementation of MIPS simplified to contain only:
  - arithmetic-logical instructions: `add, sub, and, or, slt`
  - arithmetic-logical instructions with immediates: `addi, subi, andi, ori, slt`
  - memory-reference instructions: `lw, sw`
  - control flow instructions: `beq, jr, j, jal`
- Generic Implementation:
  - use the `program counter (PC)` to supply instruction address
  - get the `instruction` from memory
  - read registers
  - use the instruction to decide exactly what to do
- Two sets of registers
  - Architecturally visible: `PC, Reg File`
  - Architecturally invisible: `Multi-cycle physical registers, Pipe registers`

# The MIPS Instruction Formats


- All MIPS instructions are 32 bits long. The three instruction formats:



# The MIPS Subset


- R-type

- *add rd, rs, rt*
- *sub, and, or, slt*
- *jr rs => Function Call / Function Call Return: jr \$ra = jr \$r31*




- LOAD and STORE

- *lw rt, rs, imm16*
- *sw rt, rs, imm16*



- BRANCH:

- *beq rs, rt, imm16*



- JUMP

- *j label*
- *jal label => Function call*

Category	Instr	Op Code	Example	Meaning
<b>Arithmetic (R format)</b>	<b>add</b>	<b>0 and 32</b>	<b>add \$s1, \$s2, \$s3</b>	<b>\$s1 = \$s2 + \$s3</b>
	<b>subtract</b>	<b>0 and 34</b>	<b>sub \$s1, \$s2, \$s3</b>	<b>\$s1 = \$s2 - \$s3</b>
<b>Data transfer (I format)</b>	<b>load word</b>	<b>35</b>	<b>lw \$s1, 100(\$s2)</b>	<b>\$s1 = Memory(\$s2+100)</b>
	<b>store word</b>	<b>43</b>	<b>sw \$s1, 100(\$s2)</b>	<b>Memory(\$s2+100) = \$s1</b>
	<b>load byte</b>	<b>32</b>	<b>lb \$s1, 101(\$s2)</b>	<b>\$s1 = Memory(\$s2+101)</b>
	<b>store byte</b>	<b>40</b>	<b>sb \$s1, 101(\$s2)</b>	<b>Memory(\$s2+101) = \$s1</b>
<b>Cond. Branch</b>	<b>br on equal</b>	<b>4</b>	<b>beq \$s1, \$s2, L</b>	<b>if (\$s1==\$s2) go to L</b>
	<b>br on not equal</b>	<b>5</b>	<b>bne \$s1, \$s2, L</b>	<b>if (\$s1 !=\$s2) go to L</b>
	<b>set on less than</b>	<b>0 and 42</b>	<b>slt \$s1, \$s2, \$s3</b>	<b>if (\$s2&lt;\$s3) \$s1=1 else \$s1=0</b>
<b>Uncond. Jump</b>	<b>jump</b>	<b>2</b>	<b>j 2500</b>	<b>go to 10000</b>
	<b>jump register</b>	<b>0 and 8</b>	<b>jr \$t1</b>	<b>go to \$t1</b>
	<b>jump and link</b>	<b>3</b>	<b>jal 2500</b>	<b>go to 10000 \$ra=PC+4</b>

\$ra is the Return Address register, which is R31

# Namesapce Data Transfer Language

- is a mechanism for describing the movement and manipulation of data between storage elements:

$R[3] \leftarrow R[5] + R[7]$

$PC \leftarrow PC + 4 + R[5]$

$R[rd] \leftarrow R[rs] + R[rt]$

$R[rt] \leftarrow \text{Mem}[R[rs] + \text{immed}]$

# First Step: Namespace Data Transfer Language Implementation (Single Cycle)

- arithmetic-logical instructions:  
`add, sub, and, or, slt`
  - `PC <- PC + 4`
  - `RegFile[rd] <- RegFile[rs] + RegFile[rt]`
- arithmetic-logical instructions with immediates:  
`addi, subi, andi, ori, slt`
  - `PC <- PC + 4`
  - `RegFile[rt] <- RegFile[rs] + sign_ext(imm)`
- memory-reference instructions: `lw, sw`
  - `PC <- PC + 4`
  - `RegFile[rt] <- Mem[rs + sign_ext(imm)]`
  - `Mem[rs + sign_ext(imm)] <= RegFile[rt]`
- control flow instructions: `beq/bne`
  - `if (rs ==/!= rt) {PC <- PC + sign_ext(imm << 2)} else {PC <- PC + 4}`
- jump instructions:
  - `jr : PC <- RegFile[rs]`
  - `j: PC <- {PC[31:28], target_address << 2}`
  - `jal`
    - `RegFile[31] <- PC + 4 // $r31 = $ra = return address register`
    - `PC <- {PC[31:28], target_address << 2}`



jump instruction: 0000 1010 1100 0101 0001 0100 0110 0010

op-code

26-bit target field from  
jump instruction

Shift Left two  
positions

32-Bit Jump Address: 0101 1011 0001 0100 0101 0001 1000 1000

High-order four bits from PC

PC: 0101 0110 0111 0110 0111 0010 1001 0100

## MIPS operands

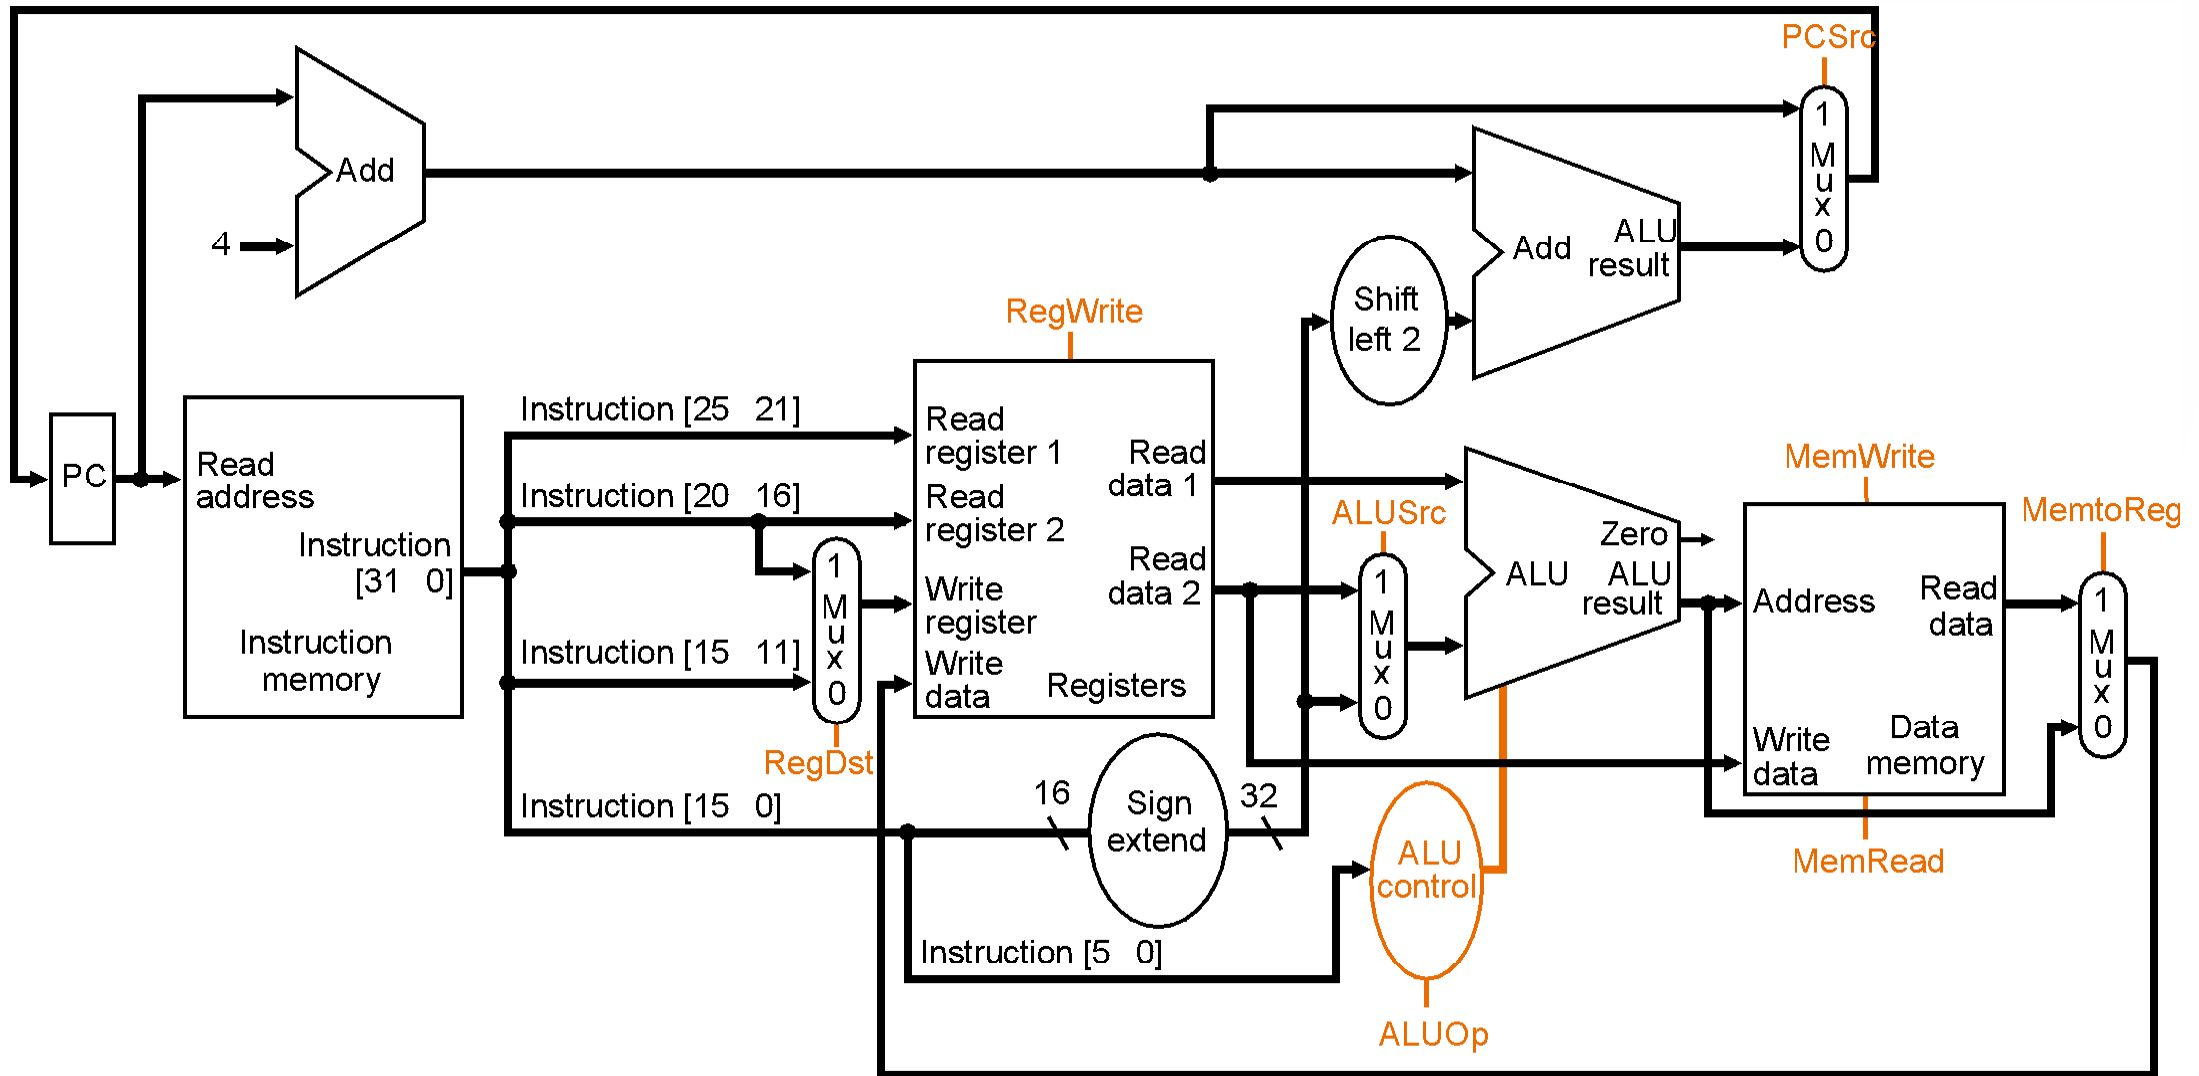
Name	Example	Comments
32 registers	\$s0-\$s7, \$t0-\$t9, \$zero, \$a0-\$a3, \$v0-\$v1, \$gp, \$fp, \$sp, \$ra, \$at	Fast locations for data. In MIPS, data must be in registers to perform arithmetic. MIPS register \$zero always equals 0. Register \$at is reserved for the assembler to handle large constants.
2 <sup>30</sup> memory words	Memory[0], Memory[4], ..., Memory[4294967292]	Accessed only by data transfer instructions. MIPS uses byte addresses, so sequential words differ by 4. Memory holds data structures, such as arrays, and spilled registers, such as those saved on procedure calls.


## MIPS assembly language


Category	Instruction	Example	Meaning	Comments
Arithmetic	add	add \$s1, \$s2, \$s3	$\$s1 = \$s2 + \$s3$	Three operands; data in registers
	subtract	sub \$s1, \$s2, \$s3	$\$s1 = \$s2 - \$s3$	Three operands; data in registers
	add immediate	addi \$s1, \$s2, 100	$\$s1 = \$s2 + 100$	Used to add constants
Data transfer	load word	lw \$s1, 100(\$s2)	$\$s1 = \text{Memory}[\$s2 + 100]$	Word from memory to register
	store word	sw \$s1, 100(\$s2)	$\text{Memory}[\$s2 + 100] = \$s1$	Word from register to memory
	load byte	lb \$s1, 100(\$s2)	$\$s1 = \text{Memory}[\$s2 + 100]$	Byte from memory to register
	store byte	sb \$s1, 100(\$s2)	$\text{Memory}[\$s2 + 100] = \$s1$	Byte from register to memory
	load upper immediate	lui \$s1, 100	$\$s1 = 100 * 2^{16}$	Loads constant in upper 16 bits
Conditional branch	branch on equal	beq \$s1, \$s2, 25	if ( $\$s1 == \$s2$ ) go to PC + 4 + 100	Equal test; PC-relative branch
	branch on not equal	bne \$s1, \$s2, 25	if ( $\$s1 \neq \$s2$ ) go to PC + 4 + 100	Not equal test; PC-relative
	set on less than	slt \$s1, \$s2, \$s3	if ( $\$s2 < \$s3$ ) $\$s1 = 1$ ; else $\$s1 = 0$	Compare less than; for beq, bne
	set less than immediate	slti \$s1, \$s2, 100	if ( $\$s2 < 100$ ) $\$s1 = 1$ ; else $\$s1 = 0$	Compare less than constant
Unconditional jump	jump	j 2500	go to 10000	Jump to target address
	jump register	jr \$ra	go to \$ra	For switch, procedure return
	jump and link	jal 2500	$\$ra = \text{PC} + 4$ ; go to 10000	For procedure call


# Putting it All Together: A Single Cycle Datapath

- We have everything except control signals

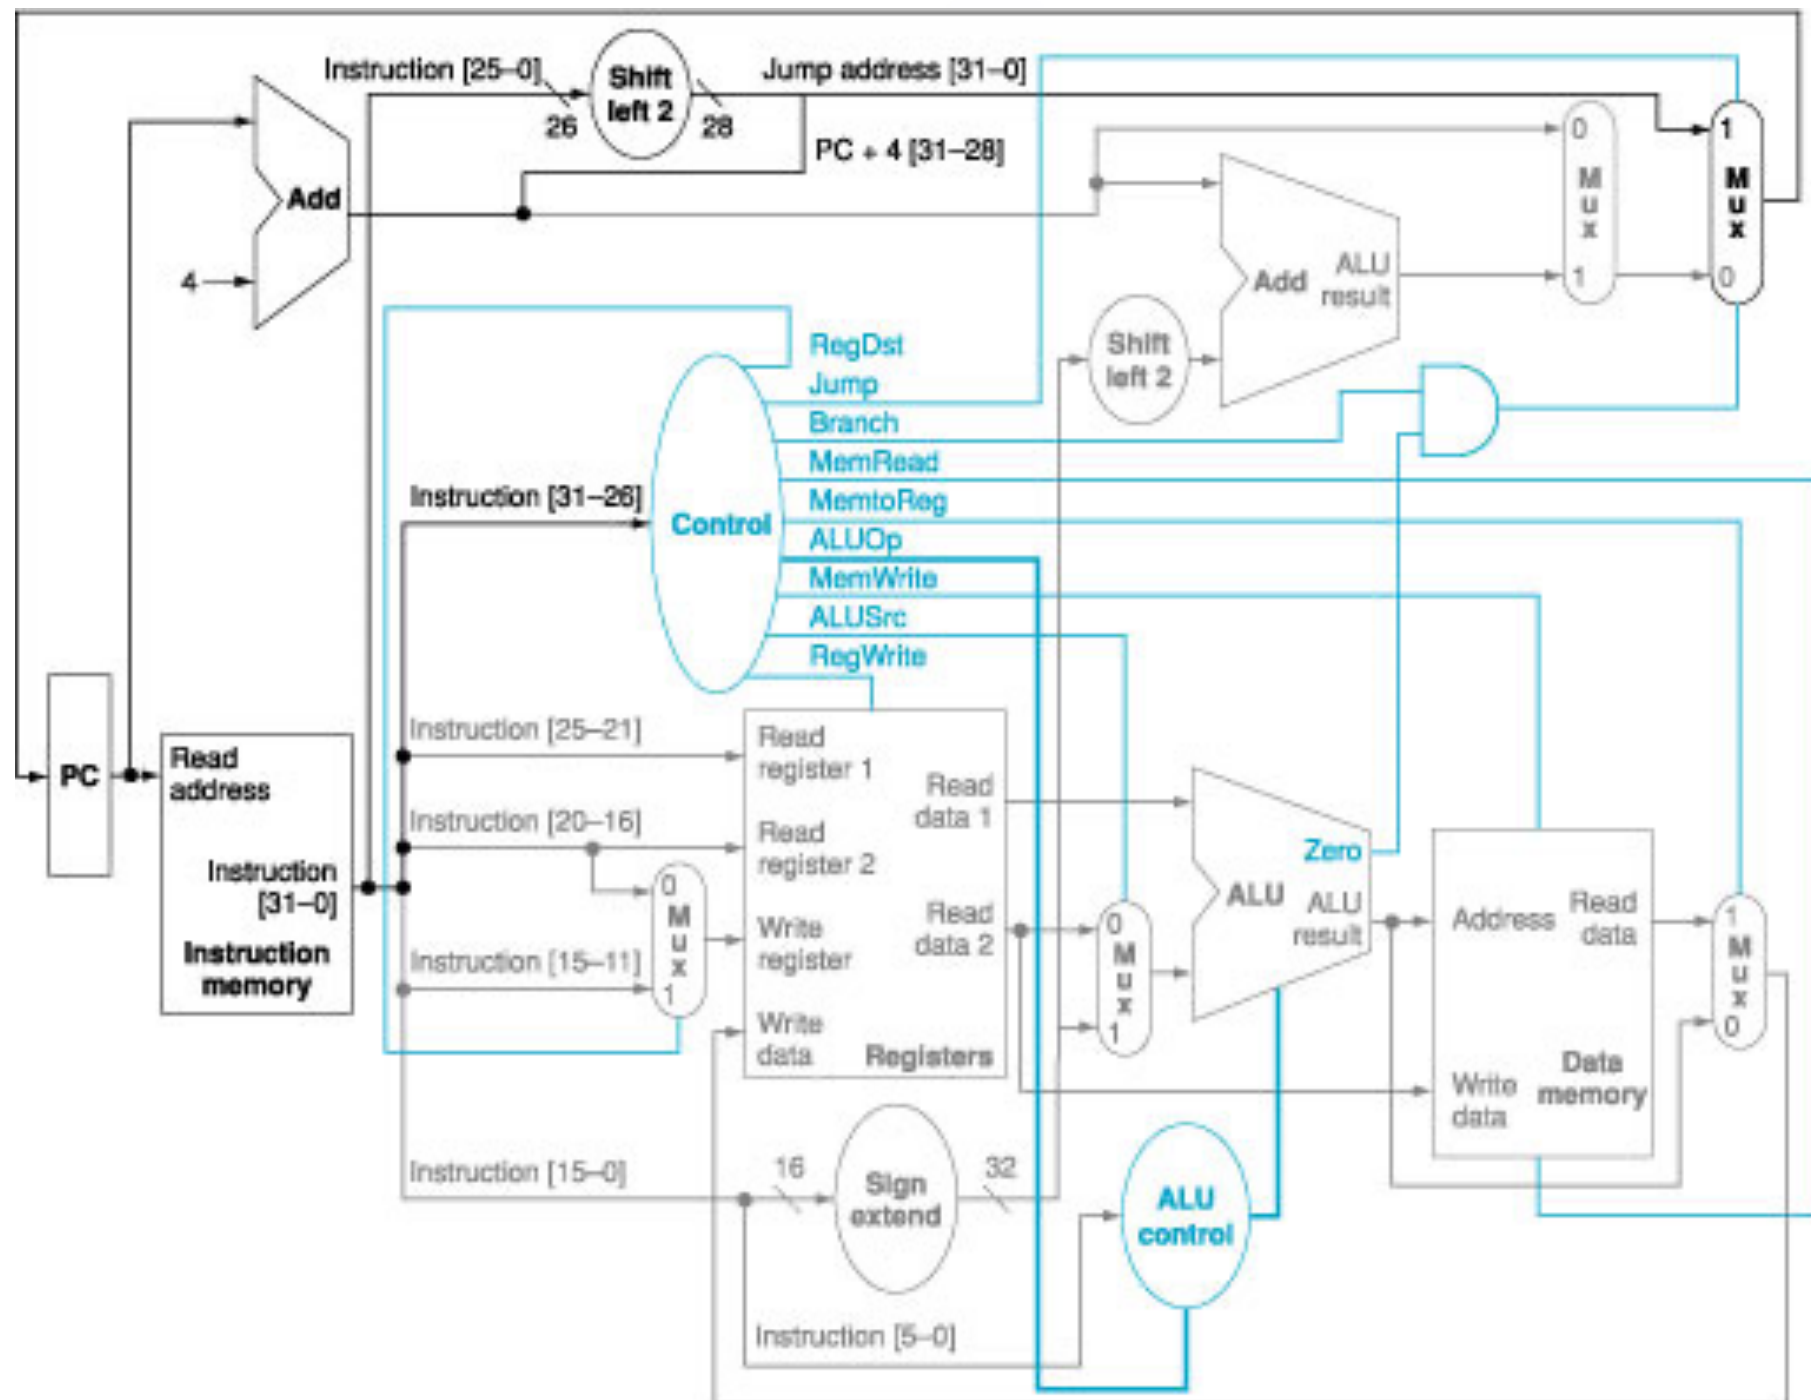


















# Single Cycle

	CPI	CT
Single Cycle		

# Why a Multiple Clock Cycle CPU?

- The problem => single-cycle cpu has a cycle time long enough to complete the **longest** instruction in the machine
- The solution => break up execution into smaller tasks, each task taking a cycle, different instructions requiring different numbers of cycles or tasks
- Other advantages => reuse of functional units (e.g., ALU, memory)
- $ET = IC * CPI * CT$



# Breaking Execution Into Clock Cycles

---

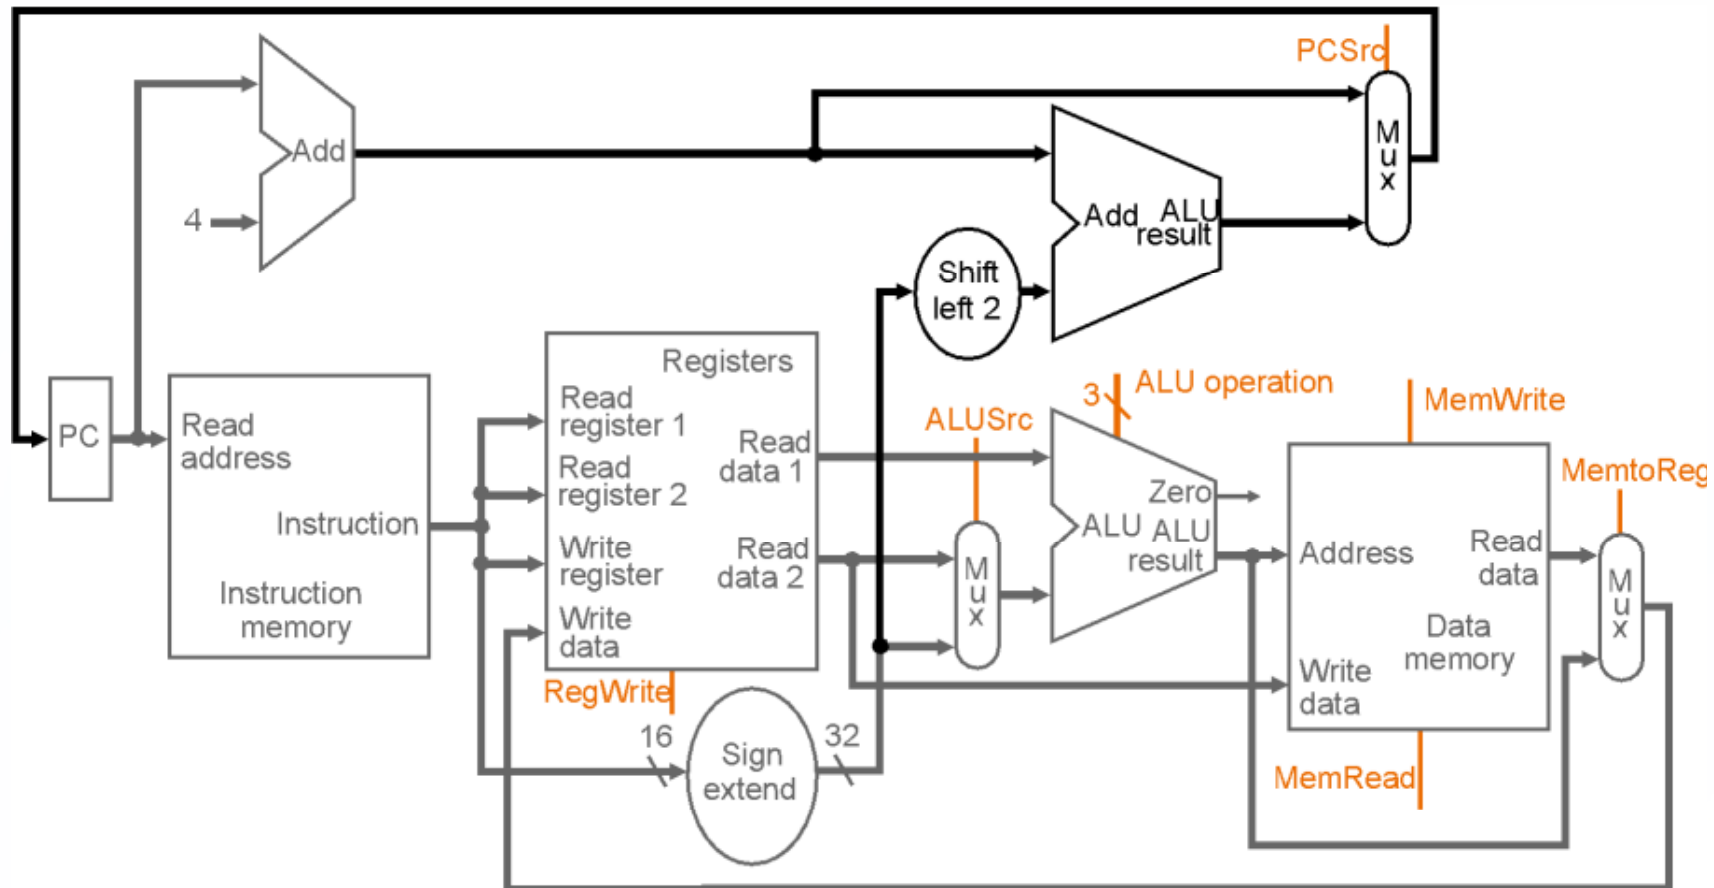
- Introduces extra registers when:
  - signal is **computed** in one clock cycle and **used** in another, AND
  - the inputs to the functional block that outputs this signal can **change** before the signal is written into a state element.
- The goal is to **balance** the amount of work done each cycle.
- Significantly complicates control. **Why?**

# Multi-cycle Review: Breaking Execution Into Clock Cycles

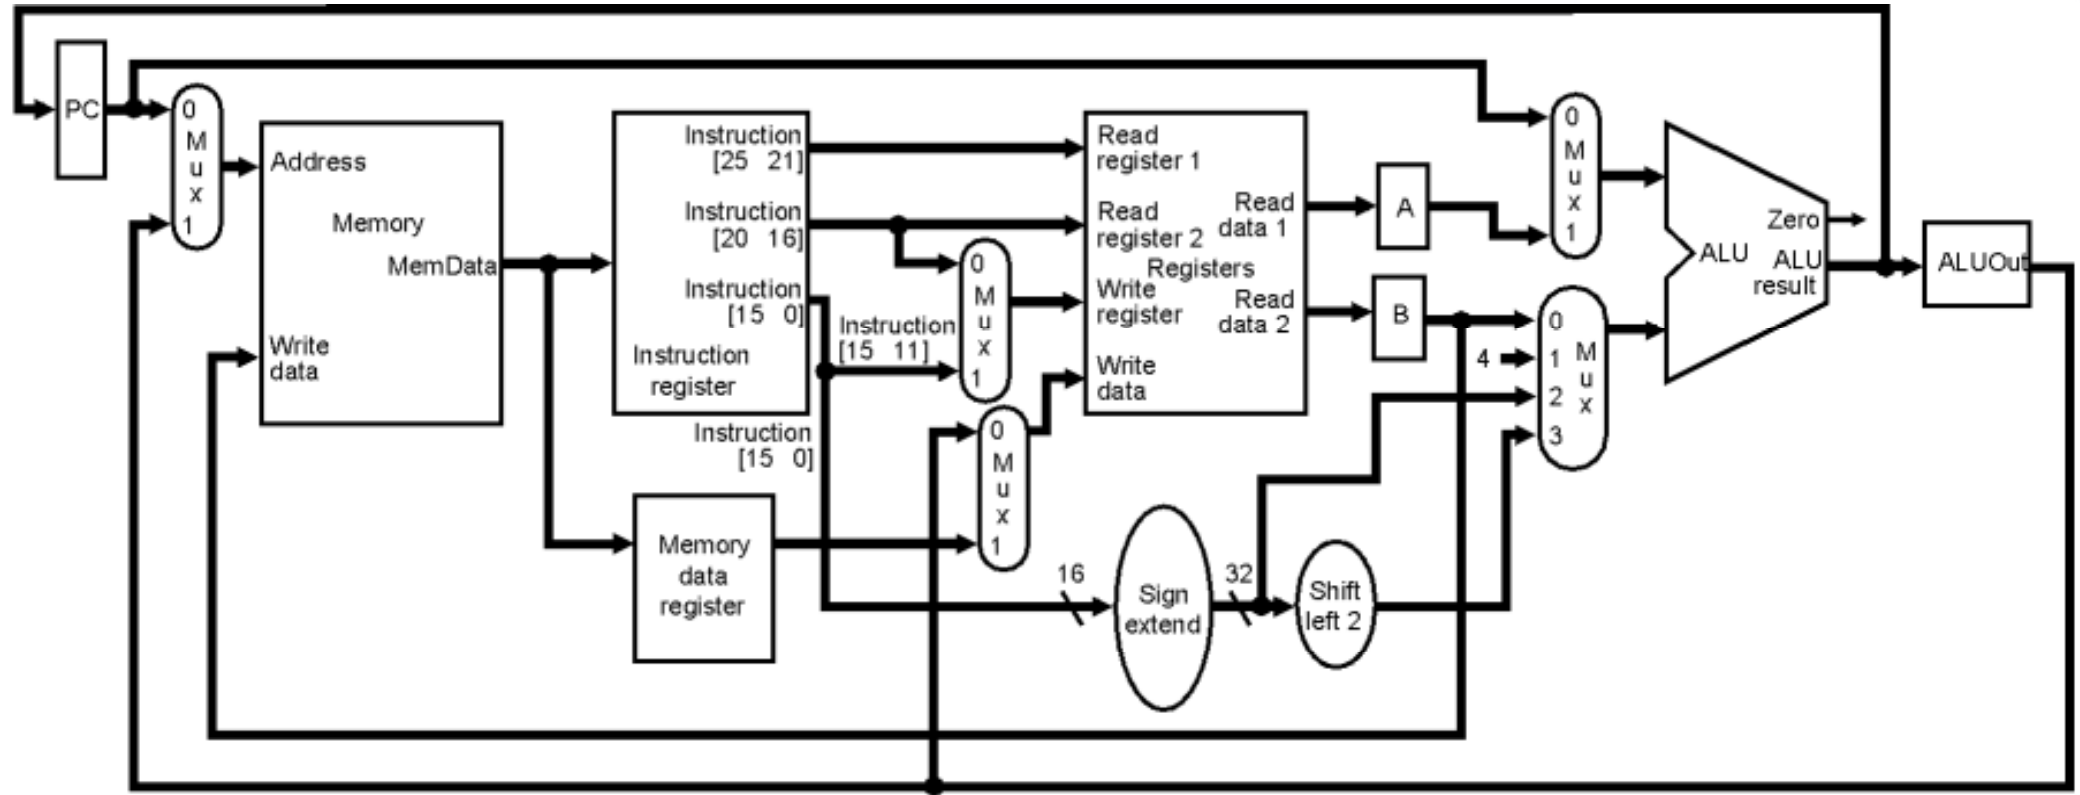
---

- We will have five execution steps (not all instructions use all five)
  - fetch
  - decode & register fetch
  - execute
  - memory access
  - write-back

# Cutting up Single Cycle



# Multicycle datapath

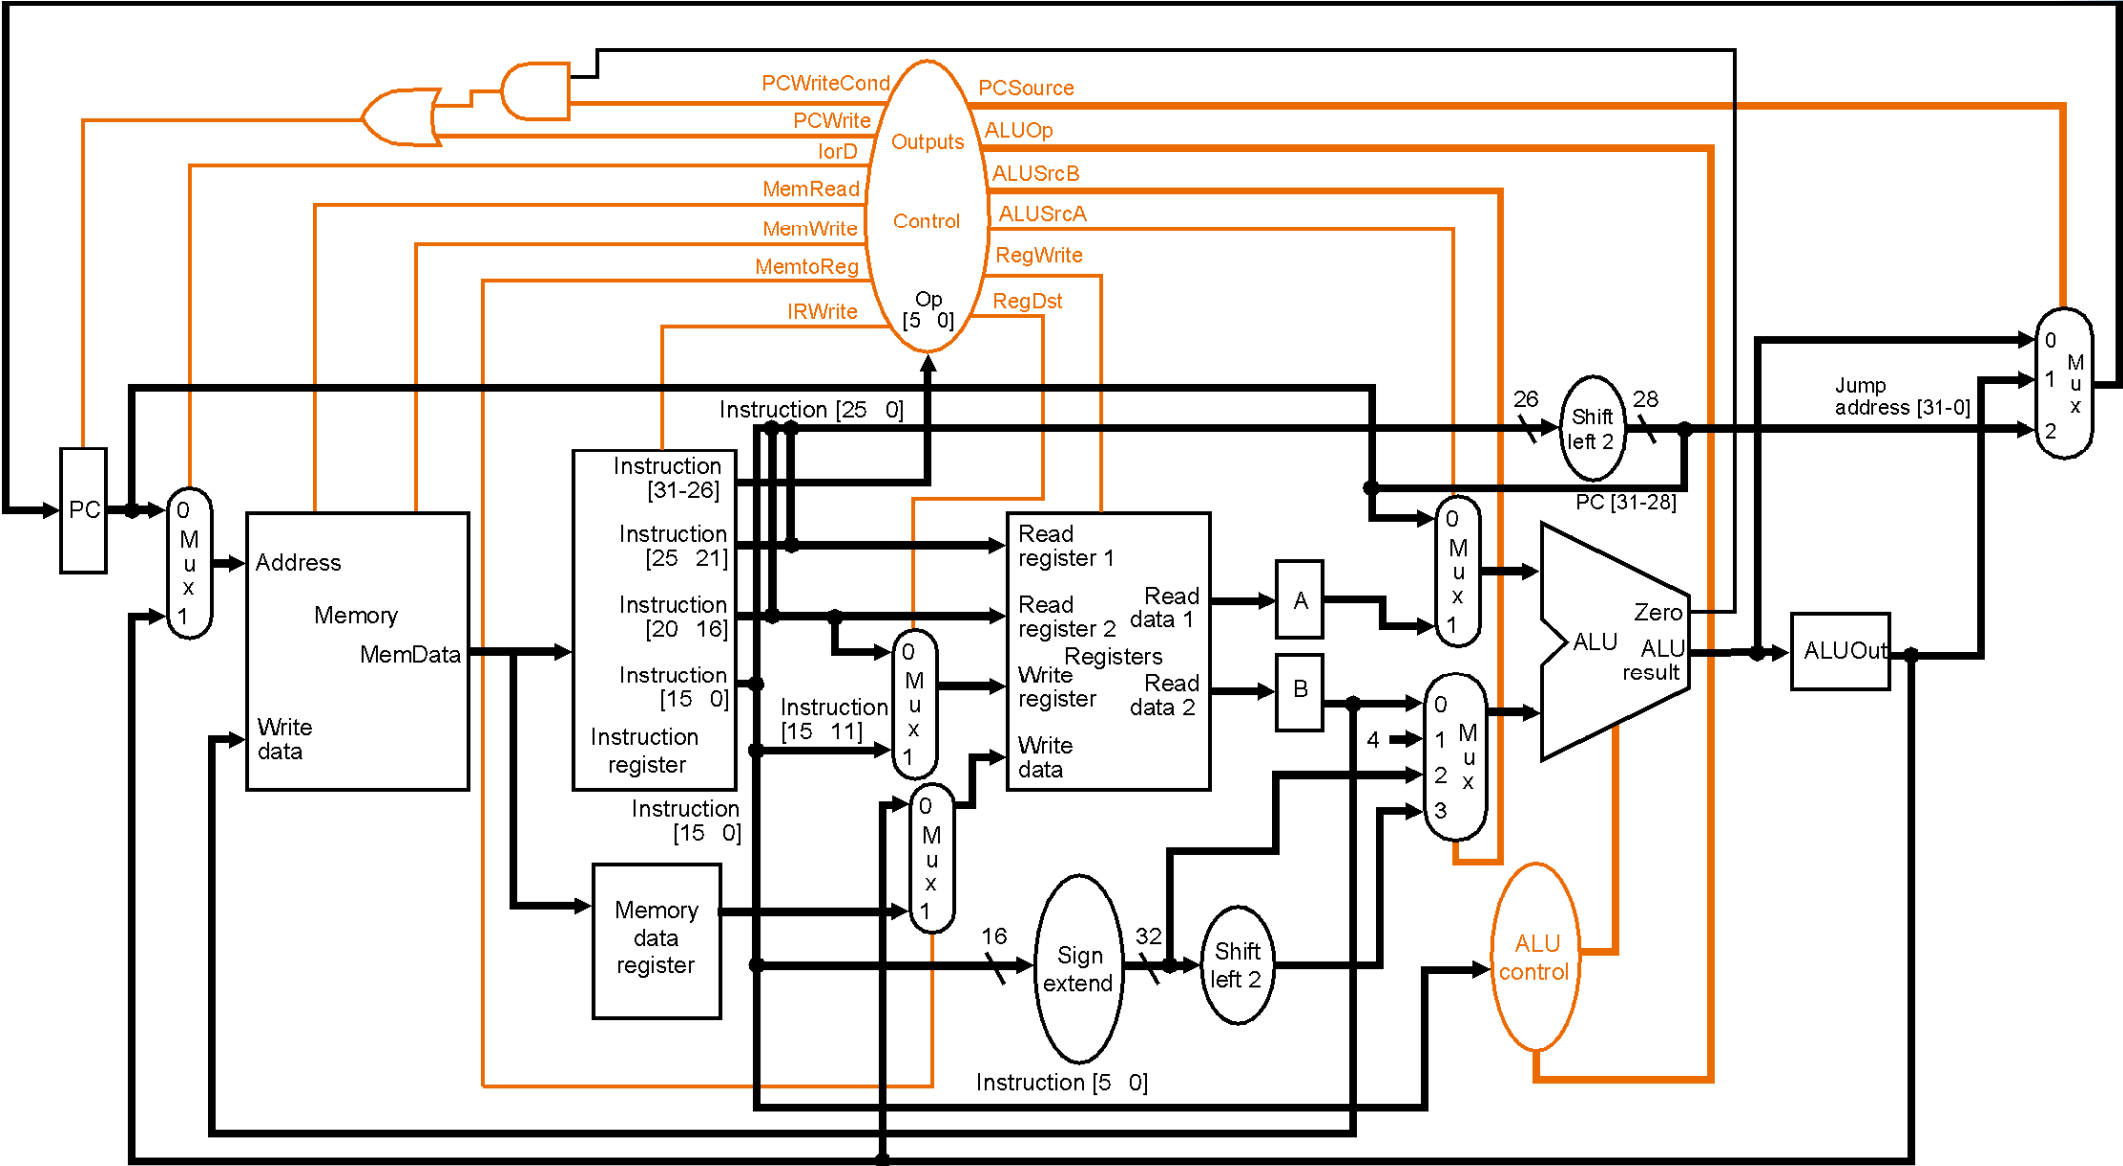


Intermediate latches.

One ALU

One memory (can we do self-modifying code?)

# Complete Multicycle Datapath

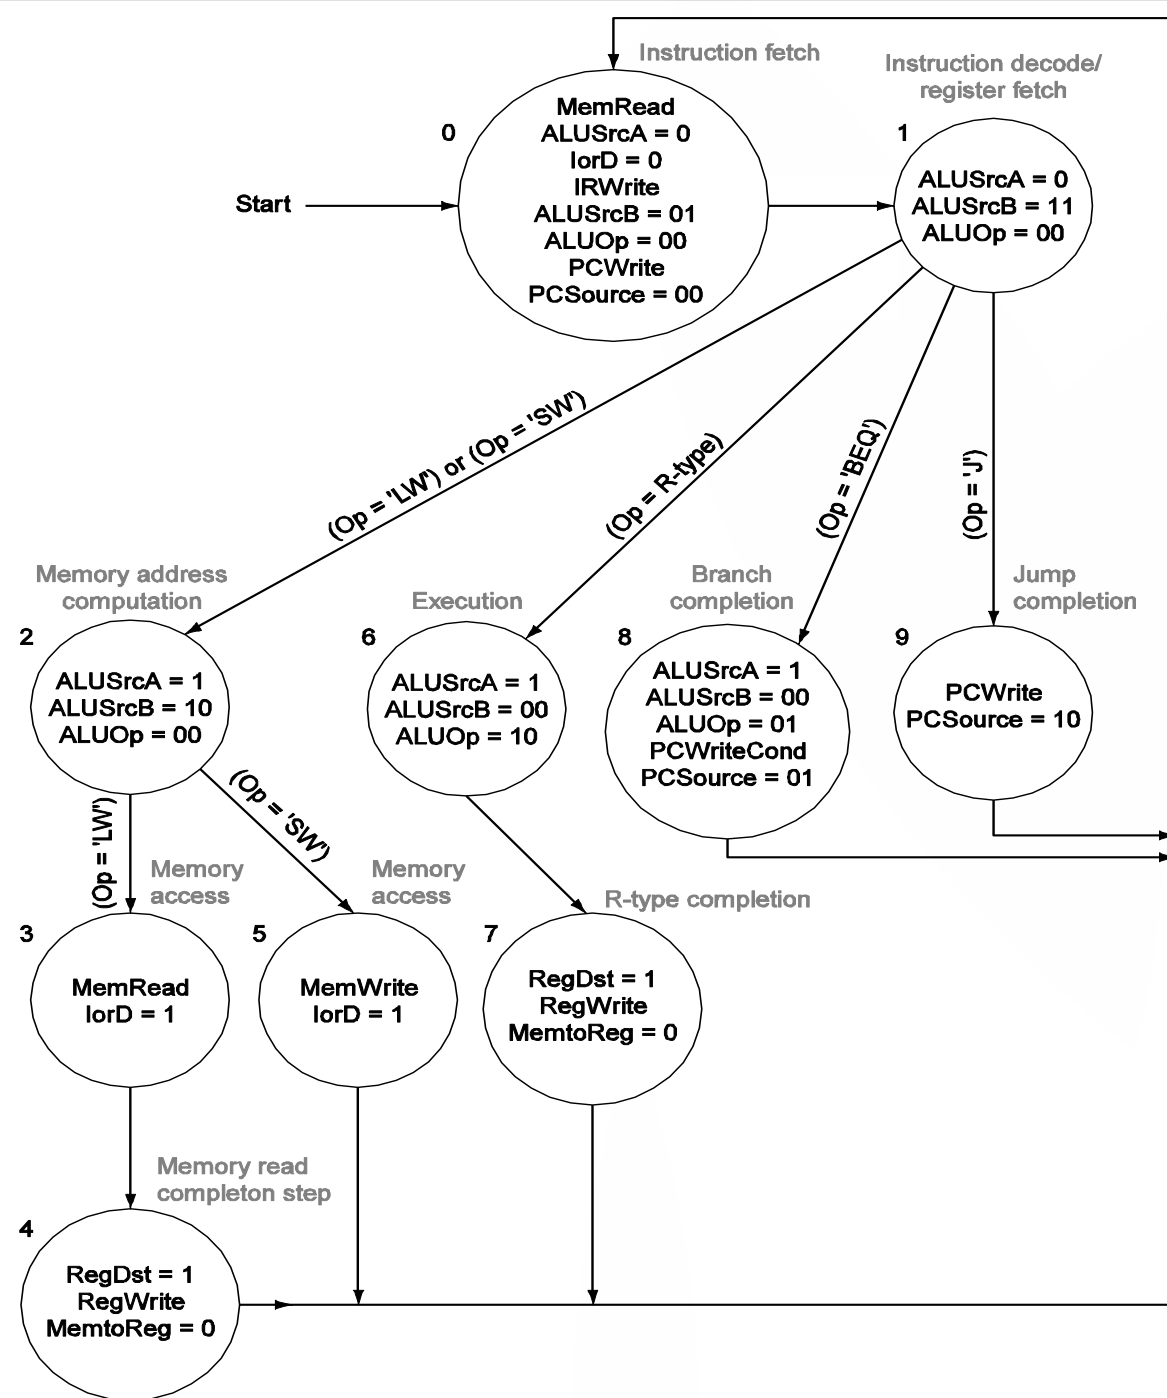


# Summary of execution steps

We can use Namespace Data Transfer Language to describe these steps

Step	R-type	Memory	Branch
Instruction Fetch	IR = Mem[PC] PC = PC + 4		
Instruction Decode/ register fetch	A = Reg[IR[25-21]] B = Reg[IR[20-16]] ALUout = PC + (sign-extend(IR[15-0]) << 2)		
Execution, address computation, branch completion	ALUout = A op B	ALUout = A + sign- extend(IR[15-0])	if (A==B) then PC=ALUout
Memory access or R- type completion	Reg[IR[15-11]] = ALUout	memory-data = Mem[ALUout] or Mem[ALUout]=B	
Write-back		Reg[IR[20-16]] = memory-data	





# Single Cycle vs. Multi-cycle

	CPI	CT
Single Cycle		
Multi-cycle		
lw		
sw		
beq		
r-type		

IF = 200ps  
ID = 200ps  
EX = 200ps  
M = 200ps  
WB = 200ps

A coworker at MIPS told you she thinks she can improve the MIPS single cycle processor performance. She tells you that she wants to remove base+displacement memory addressing and replace it with register direct addressing since only 10% of instructions are memory instructions which use a non-zero displacement.

Your coworker asks you whether or not this would be a good idea. You say?  
(Hint. Think about implications to the ISA as well as hardware)

Selection	Good idea?	Reason
A	Yes	IC stays the same. CPI stays the same. CT decreases (20%)
B	Yes	IC increases by ~10%. CPI stays the same. CT decreases (20%)
C	No	IC stays the same. CPI stays the same. CT increases (20%)
D	No	IC increases by ~10%. CPI stays the same. CT stays the same.
E	No	IC stays the same. CPI stays the same. CT stays the same. Complexity increases.



## What type of processor encouraged having many different CISC instructions?

	Processor Type	Best explanation
A	Single-Cycle	Specialized CISC instructions improved performance
B	Single-Cycle	The specialized nature of CISC instructions meant a variety of instructions were needed just to perform most tasks
C	Multi-Cycle	Specialized CISC instructions improved performance
D	Multi-Cycle	The specialized nature of CISC instructions meant a variety of instructions were needed just to perform most tasks
E	Pipeline	Specialized CISC instructions improved performance

Suppose you work on an embedded ***multi-cycle MIPS32 processor*** and your software team tells you that every program which executes has to go through memory and zero 1k bytes of data fairly often (averages 10% of ET). You realize you could just have a single instruction do this called zero1k (rs) which does:


$M[rs] = 0 \dots M[rs+1020] = 0.$

Your coworker thinks you are crazy. You reply?

Selection	Crazy?	Reason
A	Yes	The complexity of such an instruction combined with no performance gain is just silly.
B	Yes	The complexity of such an instruction combined with minimal performance gain (<5%) is not worth it.
C	No	The minimal performance gains (<5%) still rationalize this simple instruction.
D	No	The significant performance gains (>5%) clearly rationalize this complex instruction.
E	Maybe	None of the above.




# Complete Multicycle Datapath

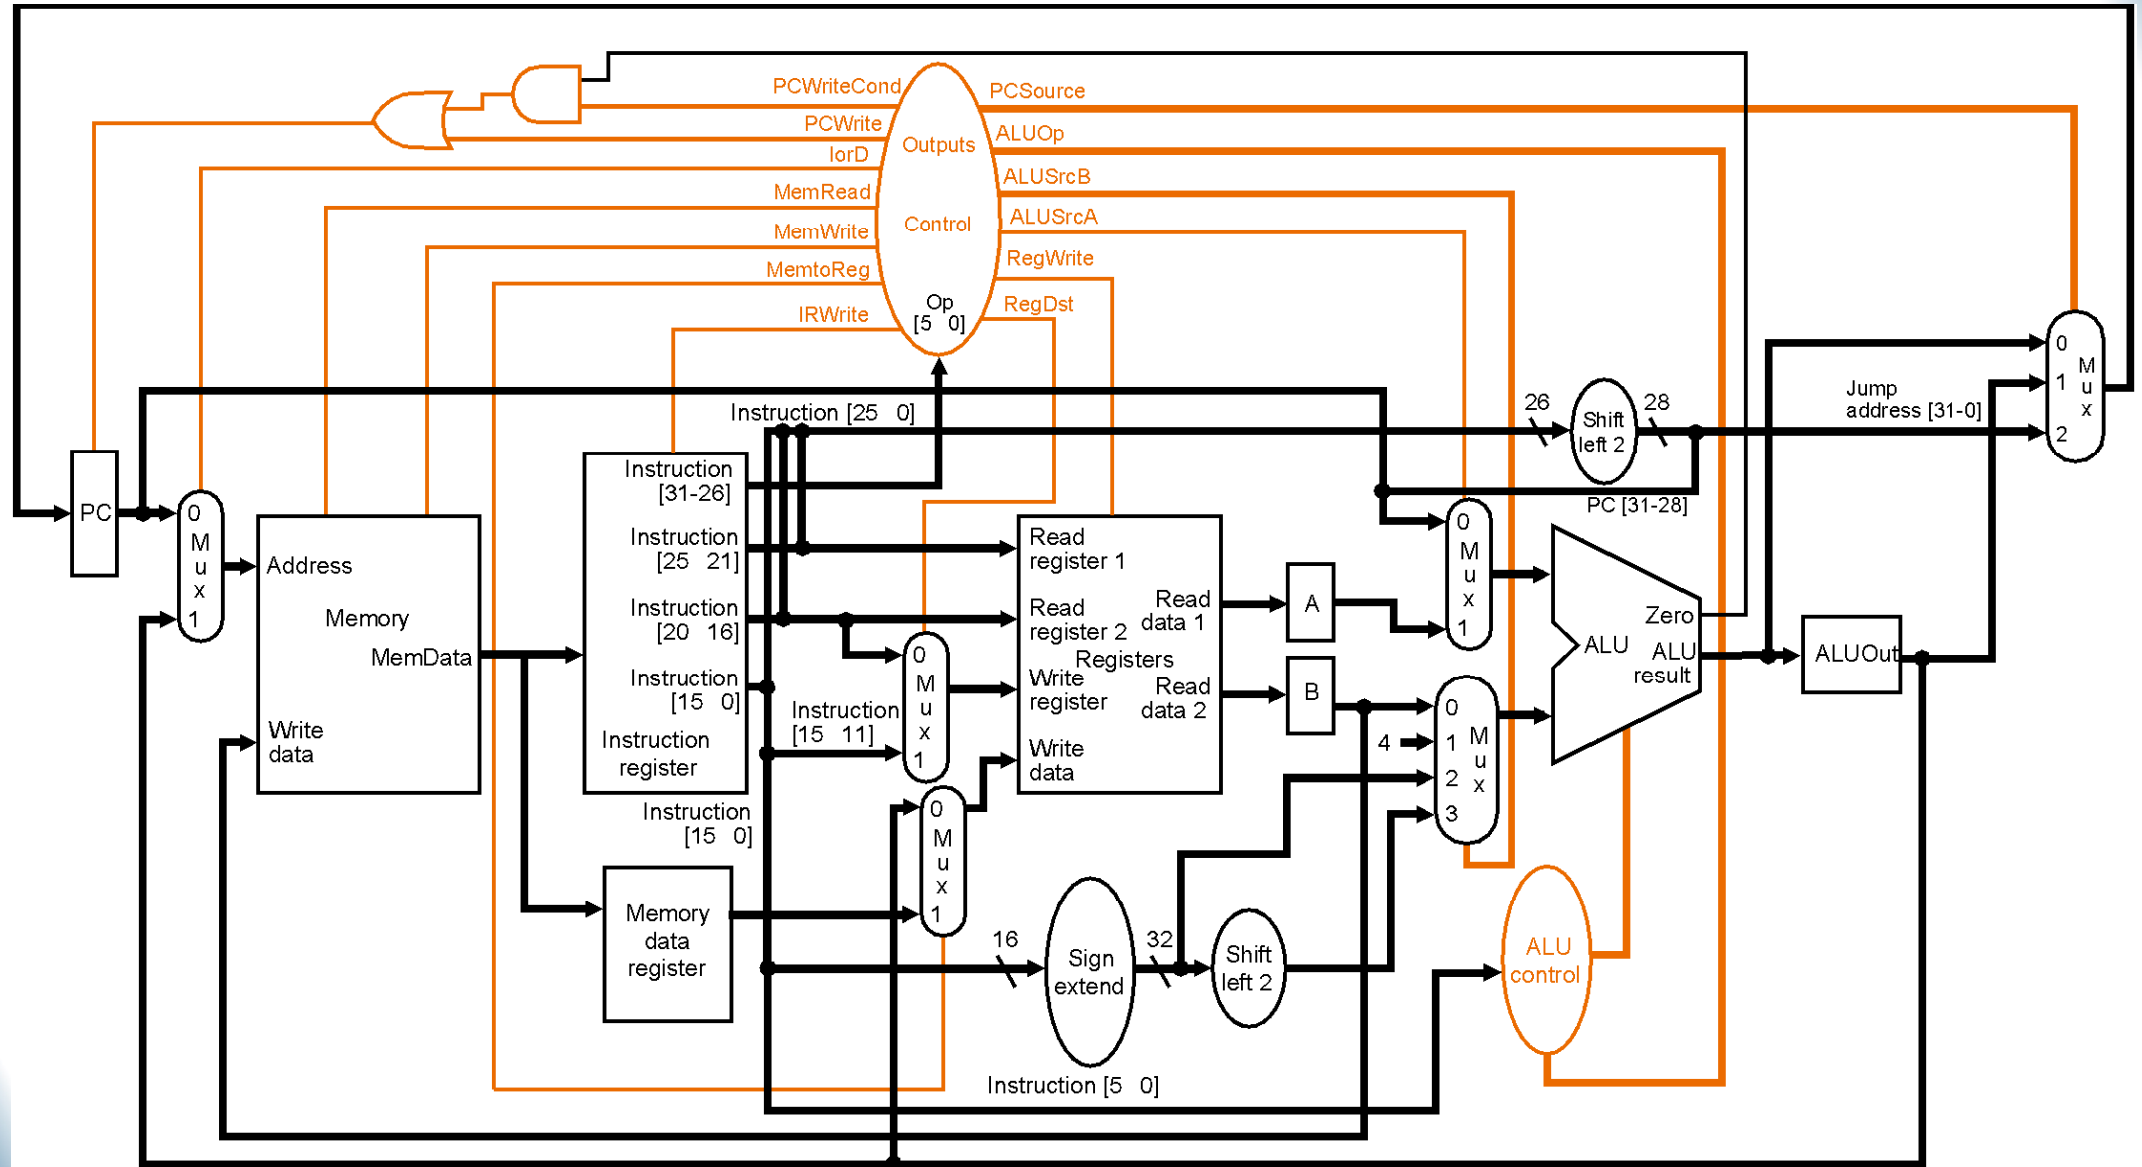




# Complete Multicycle Datapath



# Complete Multicycle Datapath

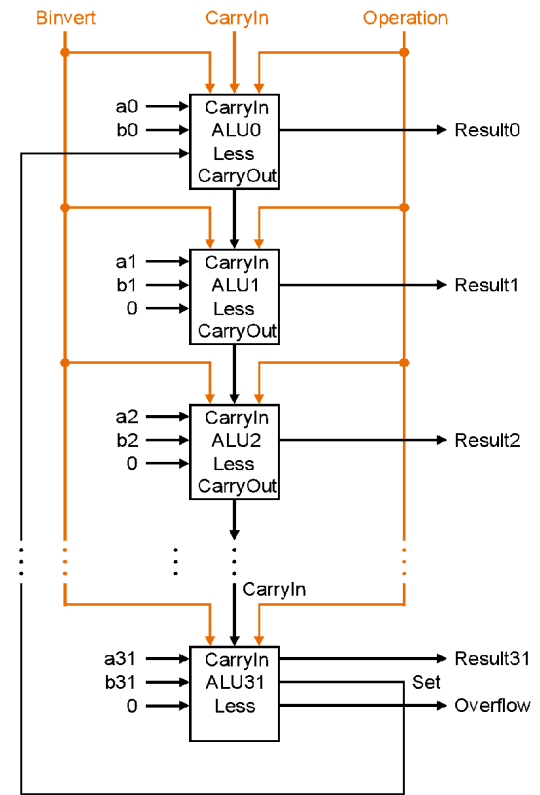




# ALU Control SIGNALS

---

# Full ALU

Consolidate to 3 wires since Binvert and CIn are always the same



what signals accomplish:  
Binvert CIn Oper  
 and?  
 or?  
 add?  
 sub?  
 beq?  
 slt?

And 0 0 0  
 Or 0 0 1  
 Add 0 0 2  
 Sub 1 1 2  
 Beq 1 1 2  
 Slt 1 1 3

# ALU control bits


- Recall: 5-function ALU

Opcode (31-26)      Function code (5-0)

ALU control input	Function	Operations
000	And	and
001	Or	or
010	Add	add, lw, sw
110	Subtract	sub, beq
111	Slt	slt


- based on Opcode (31-26) and Function code (5-0) from instruction
- ALU doesn't need to know all opcodes--we will summarize opcode with ALUOp (2 bits):

00 - lw,sw      01 - beq    10 - R-format



## Generating ALU control

Instruction opcode	ALUOp	Instruction operation	Function code	Desired ALU action	ALU control input
lw	00	load word	xxxxxx	add	010
sw	00	store word	xxxxxx	add	010
beq	01	branch eq	xxxxxx	subtract	110
R-type	10	add	100000	add	010
R-type	10	subtract	100010	subtract	110
R-type	10	AND	100100	and	000
R-type	10	OR	100101	or	001
R-type	10	slt	101010	slt	111



Essentially a truth table, and we can design logic to do this.